

param_exp_results-with_clustering_metrics-50train50run

June 20, 2023

```
[ ]: import sys
sys.path.insert(1, '/Users/madisonthantu/Desktop/DREAM/t-recs')

import trecs

import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
from scipy.ndimage import gaussian_filter1d
from collections import defaultdict
from chaney_utils import (
    load_sim_results,
    graph_relative_to_ideal,
    merge_results,
    graph_metrics,
    graph_metrics_by_axis,
    graph_relative_to_global_by_axis,
    transform_relative_to_global
)
import warnings
warnings.simplefilter("ignore")

import itertools

[ ]: DEBUG = True # don't save figures

[ ]: # results = merge_results(["param_exp_results/repeated_training",
    ↪ "param_exp_results/single_training"])

results_path = "param_exp_results/with_clustering_metrics/50train50run"

results = merge_results([results_path])
print(results.keys())
```

```

model_keys = ['0.0drift_0attention_0retraining', '0.
↳0drift_0attention_1retraining', '0.0drift_-0.8attention_0retraining', '0.
↳0drift_-0.8attention_1retraining', '0.05drift_0attention_0retraining', '0.
↳05drift_0attention_1retraining', '0.05drift_-0.8attention_0retraining', '0.
↳05drift_-0.8attention_1retraining', '0.1drift_0attention_0retraining', '0.
↳1drift_0attention_1retraining', '0.1drift_-0.8attention_0retraining', '0.
↳1drift_-0.8attention_1retraining']
print(model_keys)

id_to_readable = dict(zip(model_keys, model_keys))
print(id_to_readable)

```

```

dict_keys(['mse', 'interaction_spread', 'global_interaction_similarity',
'inter_cluster_interaction_similarity', 'intra_cluster_interaction_similarity',
'mean_global_cosine_sim', 'mean_intra_cluster_cosine_sim',
'mean_inter_cluster_cosine_sim', 'mean_cosine_sim_per_cluster',
'mean_cluster_distance_from_centroid', 'mean_global_distance_from_centroid',
'mean_distance_from_centroid_per_cluster'])
['0.0drift_0attention_0retraining', '0.0drift_0attention_1retraining',
'0.0drift_-0.8attention_0retraining', '0.0drift_-0.8attention_1retraining',
'0.05drift_0attention_0retraining', '0.05drift_0attention_1retraining',
'0.05drift_-0.8attention_0retraining', '0.05drift_-0.8attention_1retraining',
'0.1drift_0attention_0retraining', '0.1drift_0attention_1retraining',
'0.1drift_-0.8attention_0retraining', '0.1drift_-0.8attention_1retraining']
{'0.0drift_0attention_0retraining': '0.0drift_0attention_0retraining',
'0.0drift_0attention_1retraining': '0.0drift_0attention_1retraining',
'0.0drift_-0.8attention_0retraining': '0.0drift_-0.8attention_0retraining',
'0.0drift_-0.8attention_1retraining': '0.0drift_-0.8attention_1retraining',
'0.05drift_0attention_0retraining': '0.05drift_0attention_0retraining',
'0.05drift_0attention_1retraining': '0.05drift_0attention_1retraining',
'0.05drift_-0.8attention_0retraining': '0.05drift_-0.8attention_0retraining',
'0.05drift_-0.8attention_1retraining': '0.05drift_-0.8attention_1retraining',
'0.1drift_0attention_0retraining': '0.1drift_0attention_0retraining',
'0.1drift_0attention_1retraining': '0.1drift_0attention_1retraining',
'0.1drift_-0.8attention_0retraining': '0.1drift_-0.8attention_0retraining',
'0.1drift_-0.8attention_1retraining': '0.1drift_-0.8attention_1retraining'}

```

```

[ ]: hyper_params = {"drift": [0.0, 0.05, 0.1], "attention_exp": [0, -0.8],
↳ "repeated_training": [0, 1]}
models = dict([(f"{p[0]}drift_{p[1]}attention_{p[2]}retraining", p) for p in
↳ itertools.product(*hyper_params.values())])

results_df = pd.DataFrame(columns=["drift", "attention_exp",
↳ "repeated_training"] + list(results.keys()))
results_df

for params in model_keys:

```

```

df = pd.DataFrame(columns=["drift", "attention_exp", "repeated_training"] +
↳list(results.keys()))
for metric in results:
    metric_results = [vals for sim_trial in results[metric][params] for
↳vals in sim_trial]
    df[metric] = metric_results
df["drift"] = models[params][0]
df["attention_exp"] = models[params][1]
df["repeated_training"] = models[params][2]
results_df = pd.concat([results_df, df])

```

```

[ ]: print(results_df.shape)
results_df.head()
# len(models)

```

(3600, 15)

```

[ ]:
drift attention_exp repeated_training      mse  interaction_spread \
0    0.0           0           0  0.090065          -937.5
1    0.0           0           0  0.090065           -2.0
2    0.0           0           0  0.090065            1.0
3    0.0           0           0  0.090065            0.0
4    0.0           0           0  0.090065           -1.0

global_interaction_similarity  inter_cluster_interaction_similarity \
0                0.001781                0.001661
1                0.002248                0.002022
2                0.003050                0.002765
3                0.003875                0.003550
4                0.004630                0.004320

intra_cluster_interaction_similarity  mean_global_cosine_sim \
0                0.001978                0.244844
1                0.002619                0.244844
2                0.003518                0.244844
3                0.004411                0.244844
4                0.005142                0.244844

mean_intra_cluster_cosine_sim  mean_inter_cluster_cosine_sim \
0                0.318086                0.200335
1                0.318086                0.200335
2                0.318086                0.200335
3                0.318086                0.200335
4                0.318086                0.200335

mean_cosine_sim_per_cluster \
0  [0.5116548734034113, 0.9312028843954681, 0.727...

```

```

1 [0.5116548734034113, 0.9312028843954681, 0.727...
2 [0.5116548734034113, 0.9312028843954681, 0.727...
3 [0.5116548734034113, 0.9312028843954681, 0.727...
4 [0.5116548734034113, 0.9312028843954681, 0.727...

```

	mean_cluster_distance_from_centroid	mean_global_distance_from_centroid \
0	0.43184	0.687343
1	0.43184	0.687343
2	0.43184	0.687343
3	0.43184	0.687343
4	0.43184	0.687343

	mean_distance_from_centroid_per_cluster
0	[0.46841372717852575, 0.777901889986511, 0.688...
1	[0.46841372717852575, 0.777901889986511, 0.688...
2	[0.46841372717852575, 0.777901889986511, 0.688...
3	[0.46841372717852575, 0.777901889986511, 0.688...
4	[0.46841372717852575, 0.777901889986511, 0.688...

```

[ ]: y_labels = dict([
    ("mse", "mse"),
    ("interaction_spread", "interaction spread"),
    ("global_interaction_similarity", "average jacard similarity"),
    ("inter_cluster_interaction_similarity", "average jacard similarity"),
    ("intra_cluster_interaction_similarity", "average jacard similarity"),
    ("mean_global_cosine_sim", "average cosine sim"),
    ("mean_intra_cluster_cosine_sim", "average cosine sim"),
    ("mean_inter_cluster_cosine_sim", "average cosine sim"),
    ("mean_cosine_sim_per_cluster", "average cosine sim"),
    ("mean_cluster_distance_from_centroid", "distance"),
    ("mean_global_distance_from_centroid", "distance"),
    ("mean_distance_from_centroid_per_cluster", "distance"),
])

model_key_pairs = [(model_keys[i], model_keys[i+1]) for i in range(0,
    ↪len(model_keys), 2)]

single_training_keys = [key[0] for key in model_key_pairs]
repeated_training_keys = [key[1] for key in model_key_pairs]

```

1 Graphing MSE

```

[ ]: metric_key = "mse"

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

```

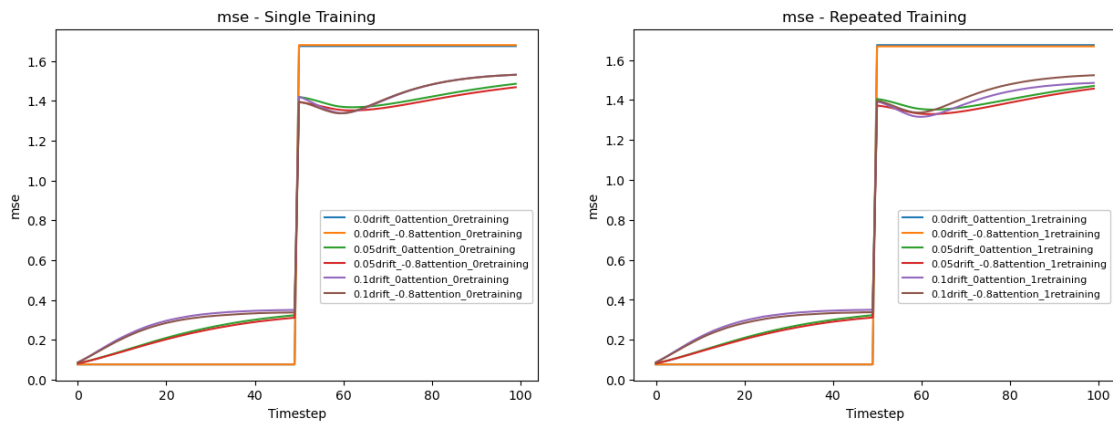
```

graph_metrics_by_axis(axes[0], results, metric_key, single_training_keys,
    ↳id_to_readable, mult_sd=0)
axes[0].set_ylabel(y_labels[metric_key])
axes[0].set_xlabel("Timestep")
axes[0].set_title(f"{metric_key} - Single Training")
axes[0].legend(facecolor='white', framealpha=1, loc='upper right',
    ↳bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axes[1], results, metric_key, repeated_training_keys,
    ↳id_to_readable, mult_sd=0)
axes[1].set_ylabel(y_labels[metric_key])
axes[1].set_xlabel("Timestep")
axes[1].set_title(f"{metric_key} - Repeated Training")
axes[1].legend(facecolor='white', framealpha=1, loc='upper right',
    ↳bbox_to_anchor=(1, 0.5), fontsize="8",)

```

[]: <matplotlib.legend.Legend at 0x1771817f0>



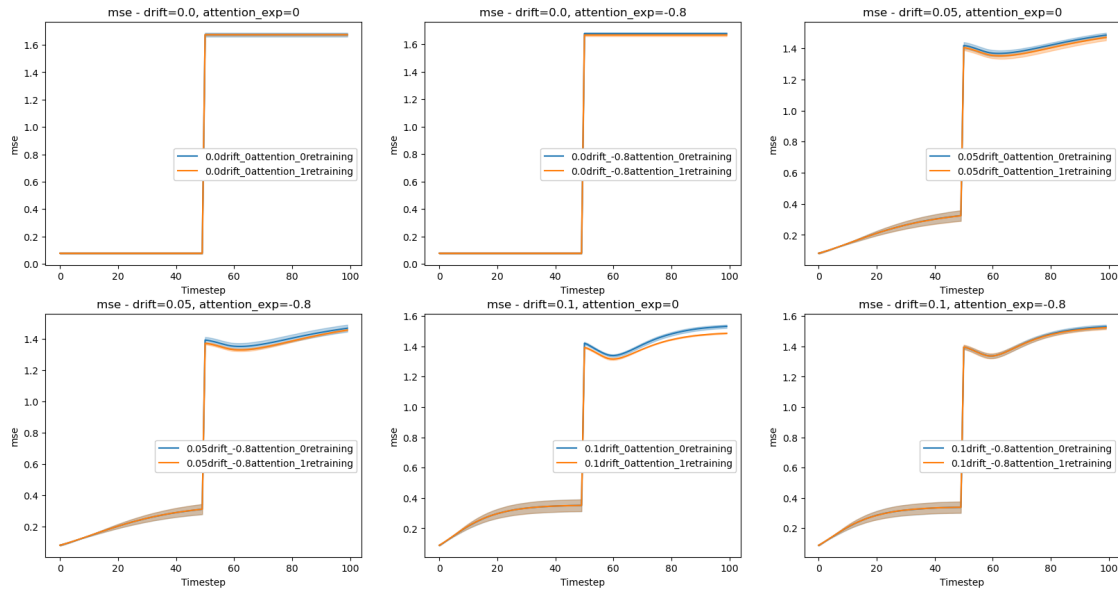
```

[ ]: metric_key = "mse"

fig, axes = plt.subplots(2, 3, figsize=(20, 10))

for i in range(len(model_key_pairs)):
    curr_ax = axes[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_key, model_key_pairs[i],
        ↳id_to_readable, mult_sd=0.5)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"{metric_key} -
        ↳drift={models[model_key_pairs[i][0]][0]},
        ↳attention_exp={models[model_key_pairs[i][0]][1]}")

```



```
[ ]: # single_training_keys = ['0.0drift_0attention_Oretraining', '0.0drift_-0.
    ↪8attention_Oretraining', '0.05drift_0attention_Oretraining', '0.05drift_-0.
    ↪8attention_Oretraining', '0.1drift_0attention_Oretraining', '0.1drift_-0.
    ↪8attention_Oretraining']
# print(single_training_keys)

# metric_key = "mse"
# # graph_relative_to_ideal(results, "mse", single_training_keys, ↪
    ↪id_to_readable, absolute_measure=False, mult_sd=1.0)
# graph_metrics(results, metric_key, single_training_keys, id_to_readable, ↪
    ↪mult_sd=0.5)
# plt.ylabel(y_labels[metric_key])
# plt.xlabel("Timestep")
# plt.legend(facecolor='white', framealpha=1, loc='upper right', ↪
    ↪bbox_to_anchor=(1.7, 1.0))
# plt.ylim(-0.1, 1.75)
# plt.xlim(0, 60)
# plt.title(f"Single Training - {metric_key}")
# if not DEBUG:
#     plt.savefig("figures/repeated_training_sim_pair.pdf", bbox_inches = ↪
    ↪"tight")

# repeated_training_keys = list(set(model_keys) - set(single_training_keys))
# print(repeated_training_keys)

# metric_key = "mse"
```

```

# # graph_relative_to_ideal(results, "mse", single_training_keys,
#   ↪id_to_readable, absolute_measure=False, mult_sd=1.0)
# graph_metrics(results, metric_key, repeated_training_keys, id_to_readable,
#   ↪mult_sd=0.5)
# plt.ylabel("Change in Jaccard Index")
# plt.xlabel("Timestep")
# plt.legend(facecolor='white', framealpha=1, loc='upper right',
#   ↪bbox_to_anchor=(1.7, 1.0))
# plt.ylim(-0.1, 1.75)
# plt.xlim(0, 60)
# plt.title(f"Repeated Training - {metric_key}")
# if not DEBUG:
#     plt.savefig("figures/repeated_training_sim_pair.pdf", bbox_inches =
#   ↪"tight")

```

2 Graphing Interaction Similarity

2.0.1 Graphing Global Interaction Similarity

```

[ ]: metric_key = "global_interaction_similarity"

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

graph_metrics_by_axis(axs[0], results, metric_key, single_training_keys,
  ↪id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right',
  ↪bbox_to_anchor=(1, 0.5), fontsize="8",)

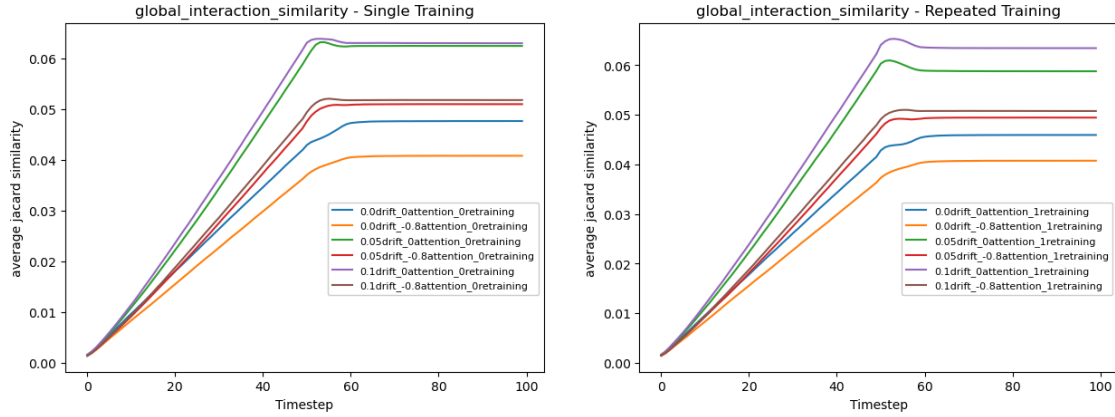
graph_metrics_by_axis(axs[1], results, metric_key, repeated_training_keys,
  ↪id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right',
  ↪bbox_to_anchor=(1, 0.5), fontsize="8",)

```

```

[ ]: <matplotlib.legend.Legend at 0x110d0f520>

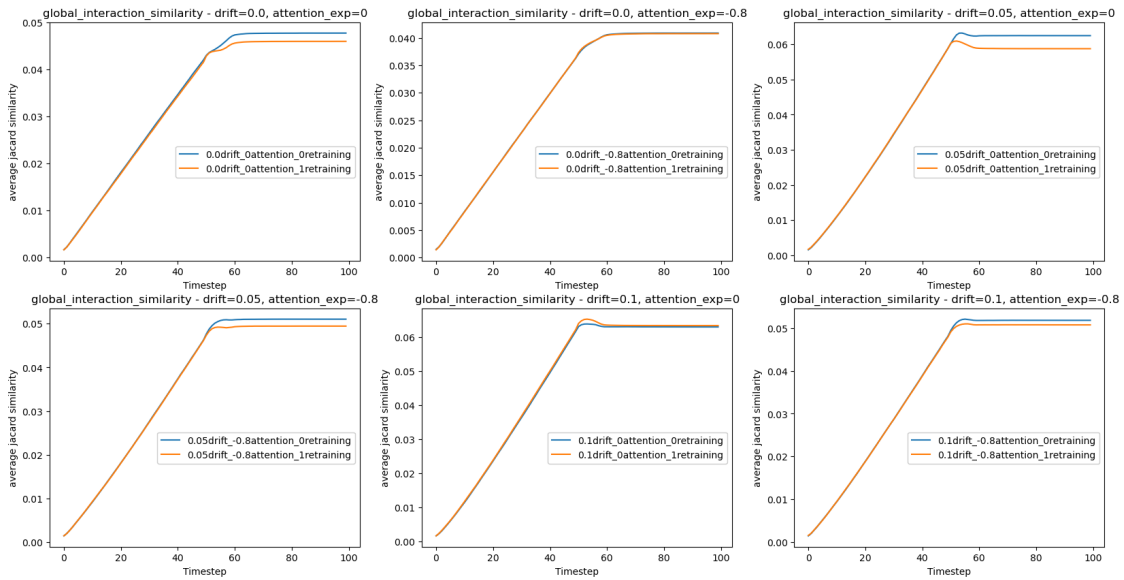
```



```
[ ]: metric_key = "global_interaction_similarity"

fig, axs = plt.subplots(2, 3, figsize=(20, 10))

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_key, model_key_pairs[i],
        ↪id_to_readable, mult_sd=0)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"{metric_key} - ↪
        ↪drift={models[model_key_pairs[i][0]][0]}, ↪
        ↪attention_exp={models[model_key_pairs[i][0]][1]}")
```



2.0.2 Graphing Inter -Cluster Interaction Similarity


```
[ ]: metric_key = "inter_cluster_interaction_similarity"

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

graph_metrics_by_axis(axs[0], results, metric_key, single_training_keys,
    ↪id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right',
    ↪bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axs[1], results, metric_key, repeated_training_keys,
    ↪id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right',
    ↪bbox_to_anchor=(1, 0.5), fontsize="8",)
```

```
[ ]: <matplotlib.legend.Legend at 0x1775060d0>
```



```
[ ]: metric_key = "inter_cluster_interaction_similarity"

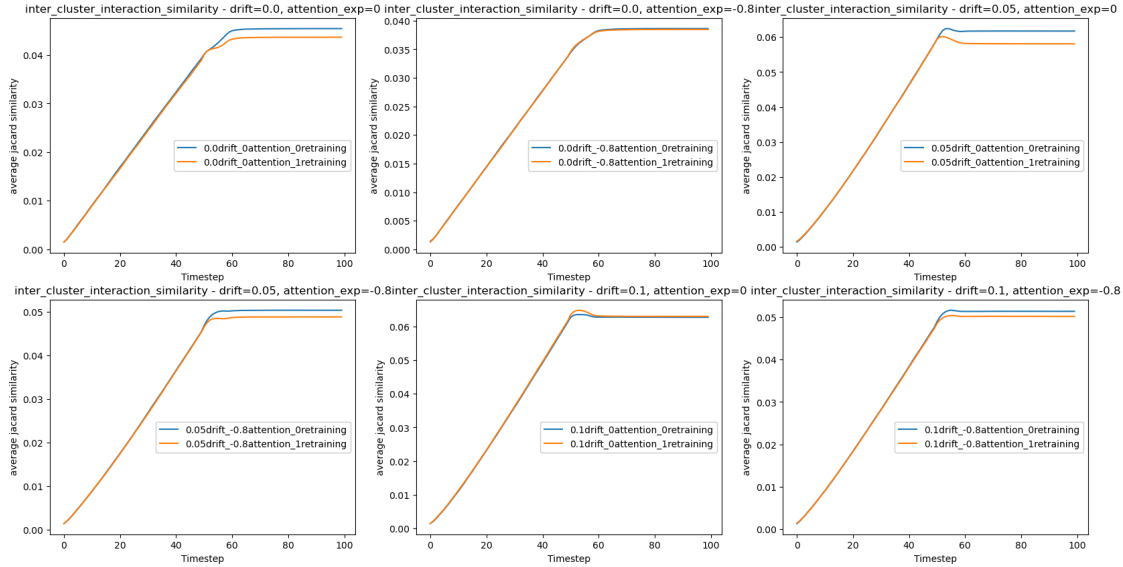
fig, axs = plt.subplots(2, 3, figsize=(20, 10))

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_key, model_key_pairs[i],
        ↪id_to_readable, mult_sd=0)
```

```

curr_ax.set_ylabel(y_labels[metric_key])
curr_ax.set_xlabel("Timestep")
curr_ax.set_title(f"{metric_key} - {training_type}")
drift={models[model_key_pairs[i][0]][0]},
attention_exp={models[model_key_pairs[i][0]][1]}")

```



2.0.3 Graphing Intra -Cluster Interaction Similarity

```

[ ]: metric_key = "intra_cluster_interaction_similarity"

single_training_keys = [key[0] for key in model_key_pairs]
repeated_training_keys = [key[1] for key in model_key_pairs]

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

graph_metrics_by_axis(axs[0], results, metric_key, single_training_keys,
    id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right',
    bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axs[1], results, metric_key, repeated_training_keys,
    id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")

```

```


axs[1].legend(facecolor='white', framealpha=1, loc='upper right',
↳bbox_to_anchor=(1, 0.5), fontsize="8",)

```

```

[ ]: <matplotlib.legend.Legend at 0x281734610>

```



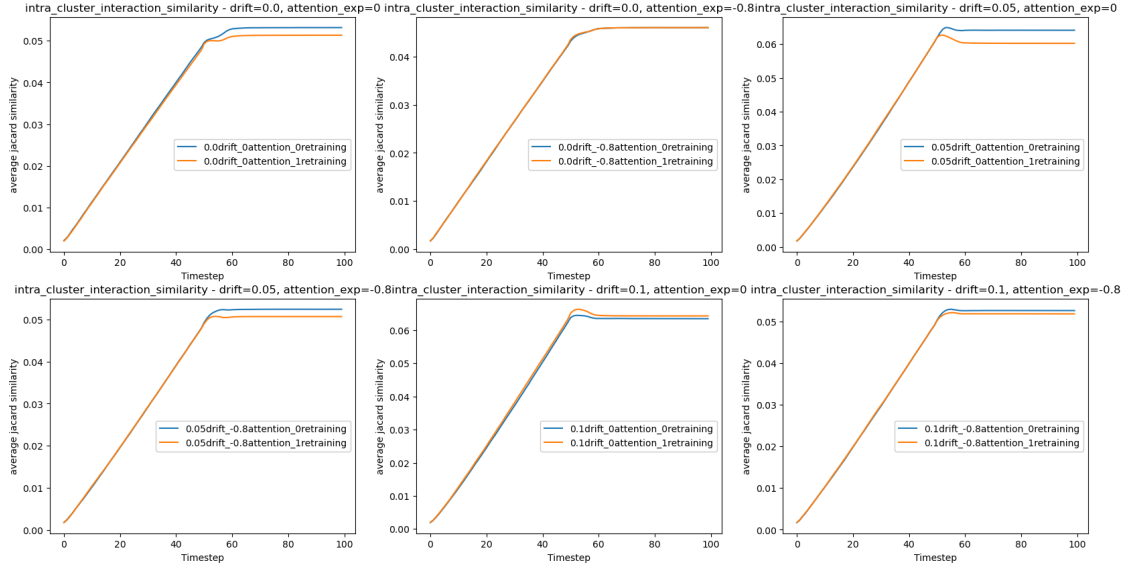
```

[ ]: metric_key = "intra_cluster_interaction_similarity"

fig, axs = plt.subplots(2, 3, figsize=(20, 10))

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_key, model_key_pairs[i],
↳id_to_readable, mult_sd=0)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"{metric_key} -
↳drift={models[model_key_pairs[i][0]][0]},
↳attention_exp={models[model_key_pairs[i][0]][1]}")

```



2.0.4 Graphing Intra -Cluster Interaction Similarity and Inter -Cluster Interaction Similarity relative to Global Interaction Similarity

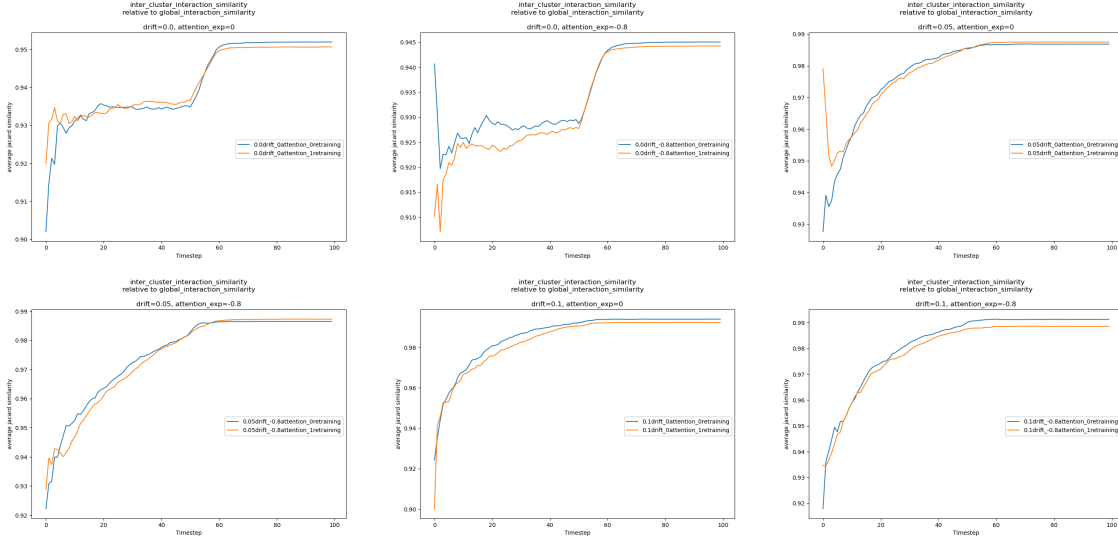
```
[ ]: global_metric_key = "global_interaction_similarity"
```

Graphing Inter -Cluster Interaction Similarity relative to Global Interaction Similarity Single training vs. Retraining

```
[ ]: metric_key = "inter_cluster_interaction_similarity"

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
    ↪metric_key, model_key_pairs[i], id_to_readable, mult_sd=0)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"{metric_key}\nrelative to
    ↪{global_metric_key}\n\ndrift={models[model_key_pairs[i][0]][0]},
    ↪attention_exp={models[model_key_pairs[i][0]][1]}")
```



```
[ ]: metric_key = "inter_cluster_interaction_similarity"


single_training_keys = [key[0] for key in model_key_pairs]
repeated_training_keys = [key[1] for key in model_key_pairs]

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

# graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
# ↪metric_key, model_key_pairs[i], id_to_readable, mult_sd=0)
graph_relative_to_global_by_axis(axs[0], results, global_metric_key,
    ↪metric_key, single_training_keys, id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right',
    ↪bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_relative_to_global_by_axis(axs[1], results, global_metric_key,
    ↪metric_key, repeated_training_keys, id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right',
    ↪bbox_to_anchor=(1, 0.5), fontsize="8",)
```

```
[ ]: <matplotlib.legend.Legend at 0x177e3c370>
```



```
[ ]: # inter_cluster_interaction_similarity = \
    ↪ results['inter_cluster_interaction_similarity']['0.1drift_-0.
    ↪ 8attention_1retraining'].mean(axis=0)
# global_interaction_similarity = results['global_interaction_similarity']['0.
    ↪ 1drift_-0.8attention_1retraining'].mean(axis=0)

# diff = np.divide(inter_cluster_interaction_similarity, \
    ↪ global_interaction_similarity)

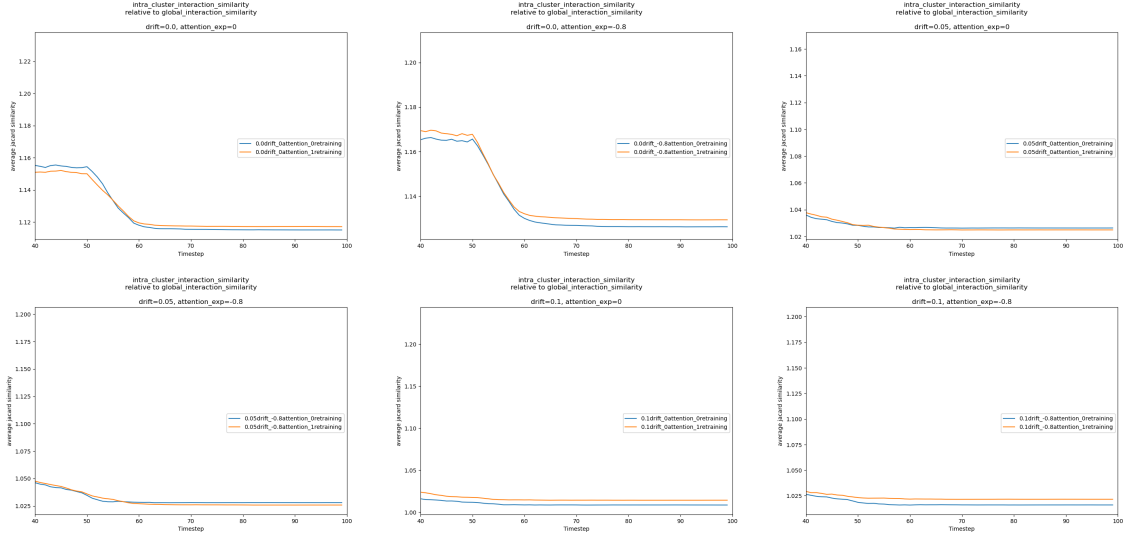
# plt.plot(diff)
```

Graphing Intra -Cluster Interaction Similarity relative to Global Interaction Similarity Single training vs. Retraining

```
[ ]: metric_key = "intra_cluster_interaction_similarity"

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, global_metric_key, \
    ↪ metric_key, model_key_pairs[i], id_to_readable, mult_sd=0)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"{metric_key}\nrelative to \
    ↪ {global_metric_key}\n\ndrift={models[model_key_pairs[i][0]][0]}, \
    ↪ attention_exp={models[model_key_pairs[i][0]][1]}")
    # curr_ax.set_ylim(-0.1, 1.75)
    curr_ax.set_xlim(40, 100)
```



```
[ ]: metric_key = "intra_cluster_interaction_similarity"

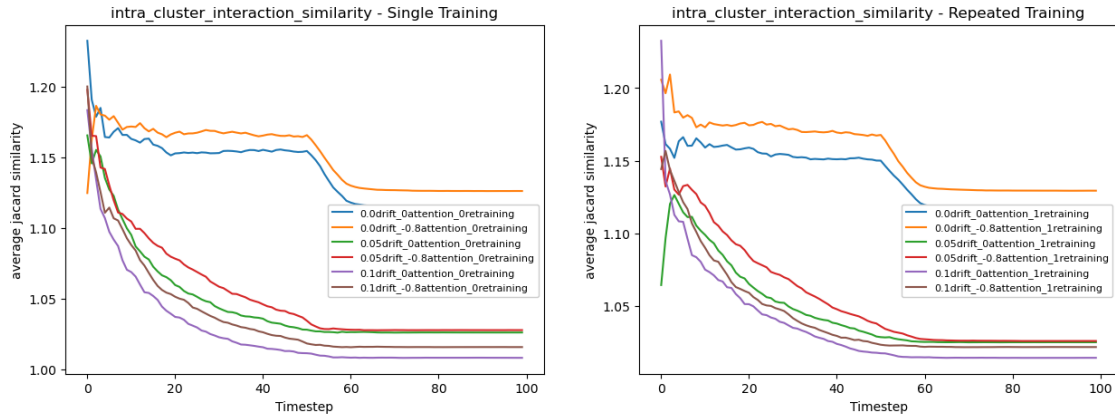
single_training_keys = [key[0] for key in model_key_pairs]
repeated_training_keys = [key[1] for key in model_key_pairs]

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

# graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
# ↪metric_key, model_key_pairs[i], id_to_readable, mult_sd=0)
graph_relative_to_global_by_axis(axs[0], results, global_metric_key,
    ↪metric_key, single_training_keys, id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right',
    ↪bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_relative_to_global_by_axis(axs[1], results, global_metric_key,
    ↪metric_key, repeated_training_keys, id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right',
    ↪bbox_to_anchor=(1, 0.5), fontsize="8",)
```

```
[ ]: <matplotlib.legend.Legend at 0x2849ade80>
```

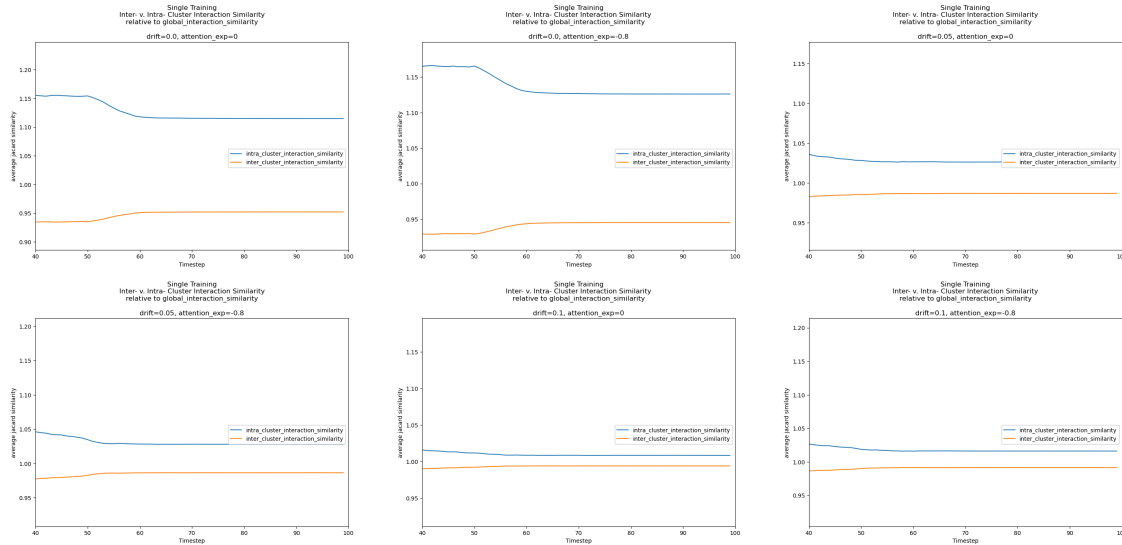


Graphing Cluster Interaction Similarity relative to Global Interaction Similarity Inter- vs. Intra-

```
[ ]: """
Single training
"""
metric_keys = ["intra_cluster_interaction_similarity",
               ↪ "inter_cluster_interaction_similarity"]

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
    ↪ metric_keys[0], (model_key_pairs[i][0]), id_to_readable, mult_sd=0,
    ↪ graph_by="metric")
    graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
    ↪ metric_keys[1], (model_key_pairs[i][0]), id_to_readable, mult_sd=0,
    ↪ graph_by="metric")
    curr_ax.set_ylabel(y_labels[metric_keys[0]])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Single Training\nInter- v. Intra- Cluster Interaction,
    ↪ Similarity\nrelative to
    ↪ {global_metric_key}\nndrift={models[model_key_pairs[i][0]][0]},
    ↪ attention_exp={models[model_key_pairs[i][0]][1]}")
    # curr_ax.set_ylim(-0.1, 1.75)
    curr_ax.set_xlim(40, 100)
```

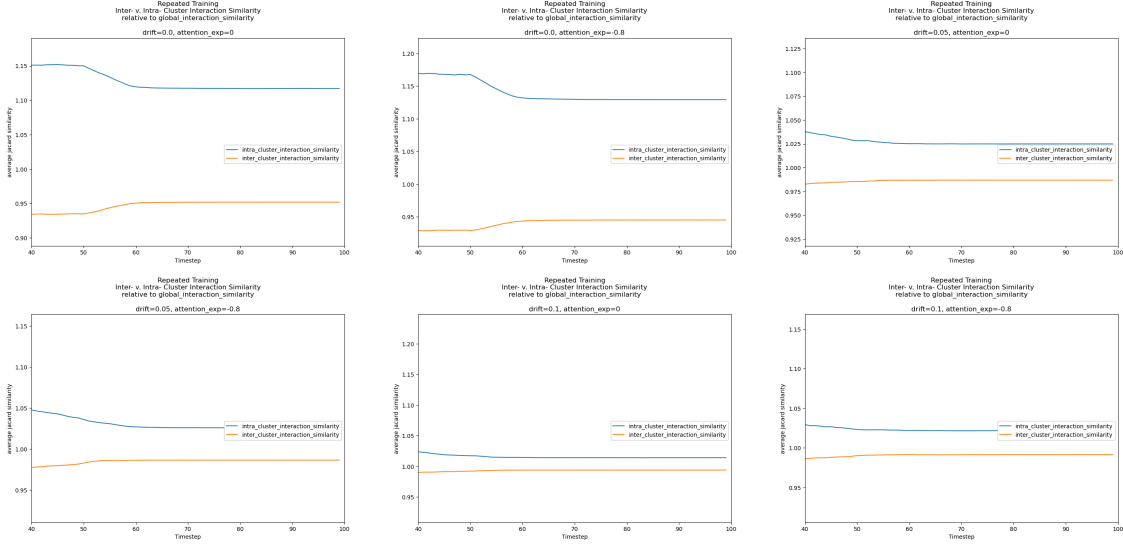



```
[ ]: """
Repeated training
"""

metric_keys = ["intra_cluster_interaction_similarity",
               ↪ "inter_cluster_interaction_similarity"]

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

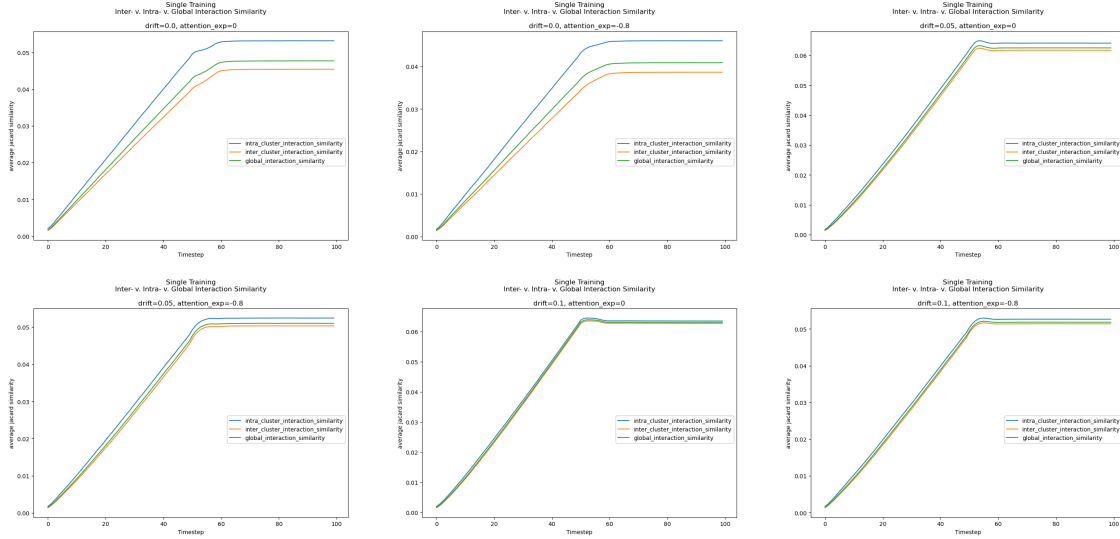
for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
    ↪ metric_keys[0], (model_key_pairs[i][1],), id_to_readable, mult_sd=0,
    ↪ graph_by="metric")
    graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
    ↪ metric_keys[1], (model_key_pairs[i][0],), id_to_readable, mult_sd=0,
    ↪ graph_by="metric")
    curr_ax.set_ylabel(y_labels[metric_keys[0]])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Repeated Training\nInter- v. Intra- Cluster Interaction_
    ↪ Similarity\nrelative to_
    ↪ {global_metric_key}\n\ndrift={models[model_key_pairs[i][0]][0]},
    ↪ attention_exp={models[model_key_pairs[i][0]][1]}")
    # curr_ax.set_ylim(-0.1, 1.75)
    curr_ax.set_xlim(40, 100)
```



```
[ ]: """
Single training
"""
metric_keys = ["intra_cluster_interaction_similarity",
               ↪ "inter_cluster_interaction_similarity", "global_interaction_similarity"]

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_keys[0],
    ↪ (model_key_pairs[i][0]), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[1],
    ↪ (model_key_pairs[i][0]), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[2],
    ↪ (model_key_pairs[i][0]), id_to_readable, mult_sd=0, graph_by="metric")
    curr_ax.set_ylabel(y_labels[metric_keys[0]])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Single Training\nInter- v. Intra- v. Global Interaction_
    ↪ Similarity\n\ndrift={models[model_key_pairs[i][0]][0]},
    ↪ attention_exp={models[model_key_pairs[i][0]][1]}")
```

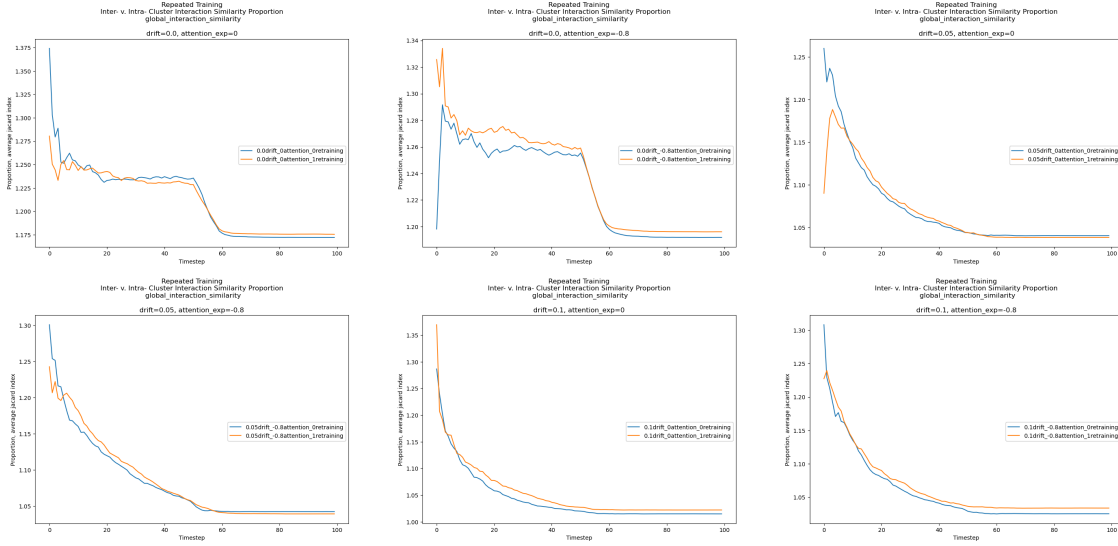


Intra-Cluster:Inter-Cluster Interaction Similarity, Single training v. Repeated training

```
[ ]: numerator = "intra_cluster_interaction_similarity"
denominator = "inter_cluster_interaction_similarity"

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, denominator, numerator,
    ↪model_key_pairs[i], id_to_readable, mult_sd=0)
    curr_ax.set_ylabel("Proportion, average jacard index")
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Repeated Training\nInter- v. Intra- Cluster Interaction_
    ↪Similarity Proportion\n_
    ↪{global_metric_key}\nndrift={models[model_key_pairs[i][0]][0]},_
    ↪attention_exp={models[model_key_pairs[i][0]][1]}")
    # curr_ax.set_ylim(-0.1, 1.75)
    # curr_ax.set_xlim(40, 100)
```



3 Graphing Cosine Similarity

3.0.1 Graphing Intra -Cluster Cosine Similarity

```
[ ]: metric_key = "mean_intra_cluster_cosine_sim"

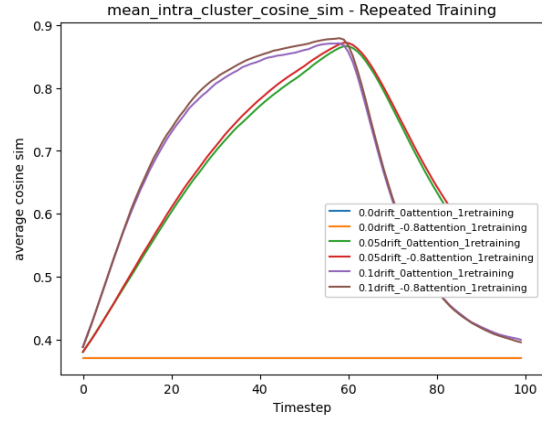

single_training_keys = [key[0] for key in model_key_pairs]
repeated_training_keys = [key[1] for key in model_key_pairs]

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

graph_metrics_by_axis(axs[0], results, metric_key, single_training_keys,
    ↳id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right',
    ↳bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axs[1], results, metric_key, repeated_training_keys,
    ↳id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right',
    ↳bbox_to_anchor=(1, 0.5), fontsize="8",)
```







```
[ ]: <matplotlib.legend.Legend at 0x28c40bb20>
```



```
[ ]: metric_key = "mean_intra_cluster_cosine_sim"

fig, axs = plt.subplots(2, 3, figsize=(20, 10))

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_key, model_key_pairs[i],
        ↪id_to_readable, mult_sd=0)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"{metric_key} - ↪
        ↪drift={models[model_key_pairs[i][0]][0]}, ↪
        ↪attention_exp={models[model_key_pairs[i][0]][1]}")
```



3.0.2 Graphing Inter -Cluster Cosine Similarity

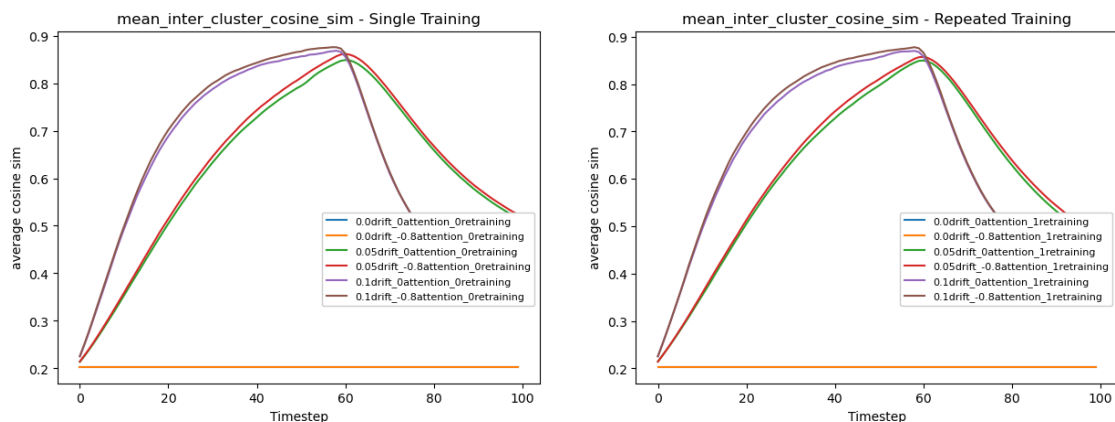
```
[ ]: metric_key = "mean_inter_cluster_cosine_sim"

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

graph_metrics_by_axis(axs[0], results, metric_key, single_training_keys,
    ↪id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right',
    ↪bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axs[1], results, metric_key, repeated_training_keys,
    ↪id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right',
    ↪bbox_to_anchor=(1, 0.5), fontsize="8",)
```

```
[ ]: <matplotlib.legend.Legend at 0x28c28caf0>
```



```
[ ]: metric_key = "mean_inter_cluster_cosine_sim"

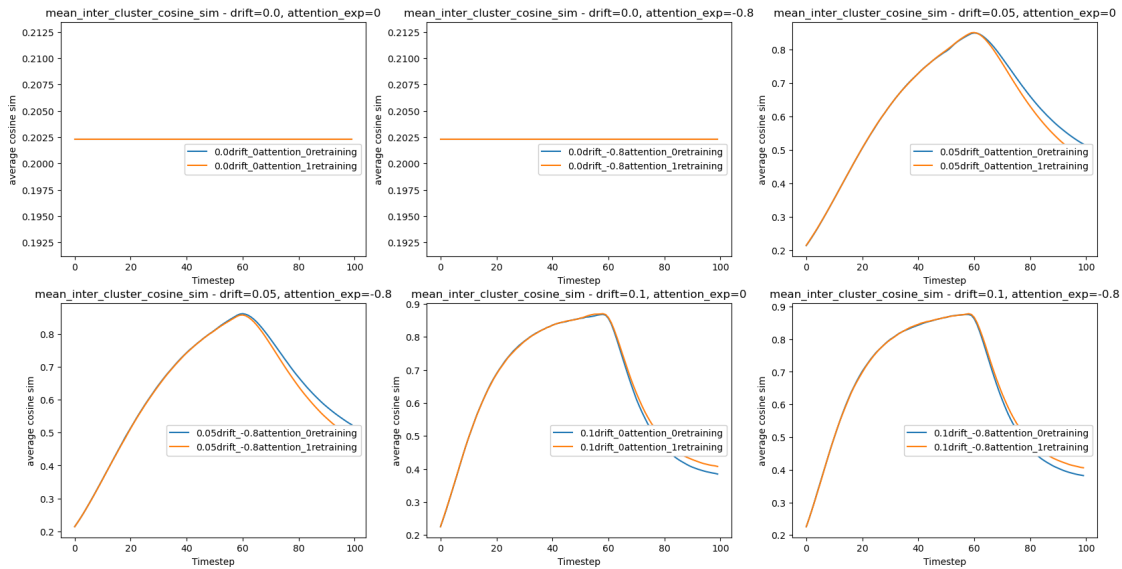
fig, axs = plt.subplots(2, 3, figsize=(20, 10))

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_key, model_key_pairs[i],
    ↪id_to_readable, mult_sd=0)
```

```

curr_ax.set_ylabel(y_labels[metric_key])
curr_ax.set_xlabel("Timestep")
curr_ax.set_title(f"{metric_key} ->
↳drift={models[model_key_pairs[i][0]][0]},<
↳attention_exp={models[model_key_pairs[i][0]][1]}")

```



Graphing Cluster Cosine Similarity relative to Global Cosine Similarity

```

[ ]: """
Single training
"""

metric_keys = ["mean_intra_cluster_cosine_sim",<
↳"mean_inter_cluster_cosine_sim", "mean_global_cosine_sim"]

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    print((model_key_pairs[i][0],))
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_keys[0],<
↳(model_key_pairs[i][0],), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[1],<
↳(model_key_pairs[i][0],), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[2],<
↳(model_key_pairs[i][0],), id_to_readable, mult_sd=0, graph_by="metric")

    curr_ax.set_xlim(75, 100)

```

```

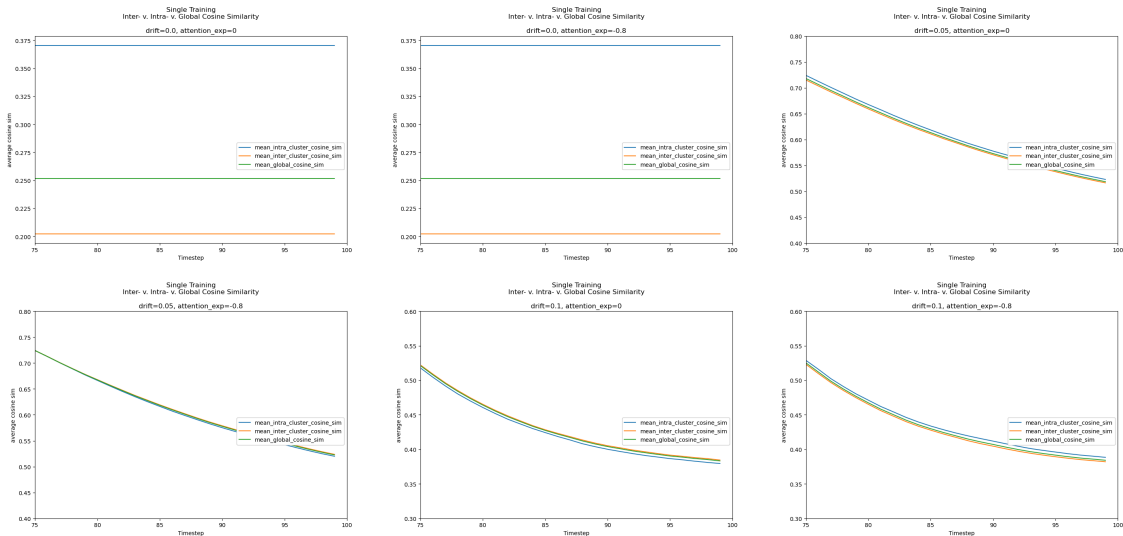
if i > 3:
    curr_ax.set_ylim(0.3, 0.6)
elif i > 1:
    curr_ax.set_ylim(0.4, 0.8)
# if i > 1:
#     curr_ax.set_ylim(0.3, 0.8)
curr_ax.set_ylabel(y_labels[metric_keys[0]])
curr_ax.set_xlabel("Timestep")
curr_ax.set_title(f"Single Training\nInter- v. Intra- v. Global Cosine_
↪Similarity\n\ndrift={models[model_key_pairs[i][0]][0]},_
↪attention_exp={models[model_key_pairs[i][0]][1]}")

```

```

('0.0drift_0attention_0retraining',)
('0.0drift_-0.8attention_0retraining',)
('0.05drift_0attention_0retraining',)
('0.05drift_-0.8attention_0retraining',)
('0.1drift_0attention_0retraining',)
('0.1drift_-0.8attention_0retraining',)

```



```

[ ]: """
Repeated training
"""
metric_keys = ["mean_intra_cluster_cosine_sim",_
↪"mean_inter_cluster_cosine_sim", "mean_global_cosine_sim"]

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):

```



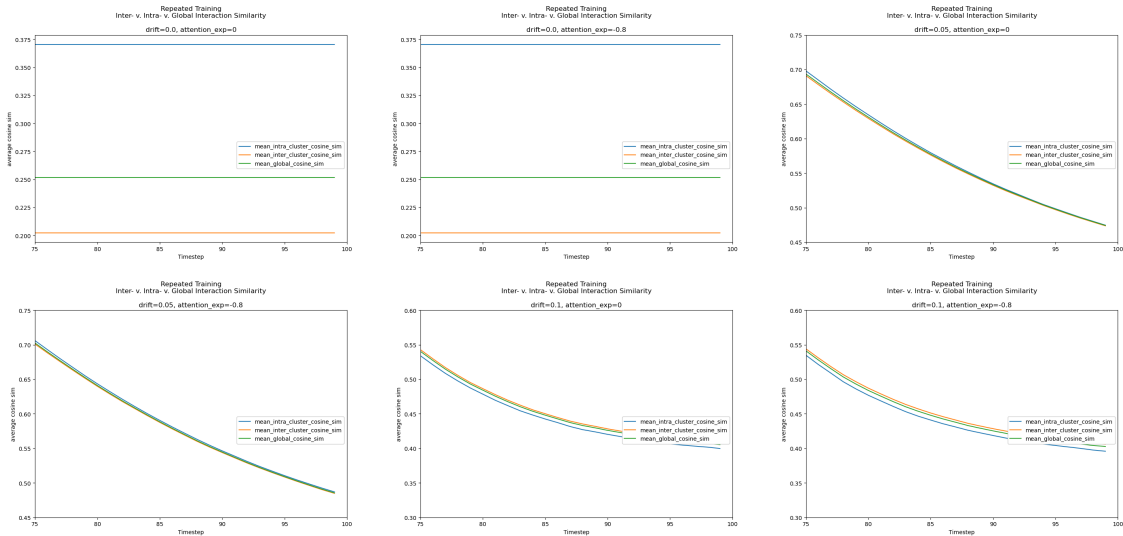
```

print(model_key_pairs[i][1])
curr_ax = axs[int(i >= 3), i%3]
graph_metrics_by_axis(curr_ax, results, metric_keys[0],
↳(model_key_pairs[i][1]), id_to_readable, mult_sd=0, graph_by="metric")
graph_metrics_by_axis(curr_ax, results, metric_keys[1],
↳(model_key_pairs[i][1]), id_to_readable, mult_sd=0, graph_by="metric")
graph_metrics_by_axis(curr_ax, results, metric_keys[2],
↳(model_key_pairs[i][1]), id_to_readable, mult_sd=0, graph_by="metric")

curr_ax.set_xlim(75, 100)
if i > 3:
    curr_ax.set_ylim(0.3, 0.6)
elif i > 1:
    curr_ax.set_ylim(0.45, 0.75)
# if i > 1:
#     curr_ax.set_ylim(0.3, 0.8)
curr_ax.set_ylabel(y_labels[metric_keys[0]])
curr_ax.set_xlabel("Timestep")
curr_ax.set_title(f"Repeated Training\nInter- v. Intra- v. Global
↳Interaction Similarity\n\ndrift={models[model_key_pairs[i][0]][0]},
↳attention_exp={models[model_key_pairs[i][0]][1]}")

```

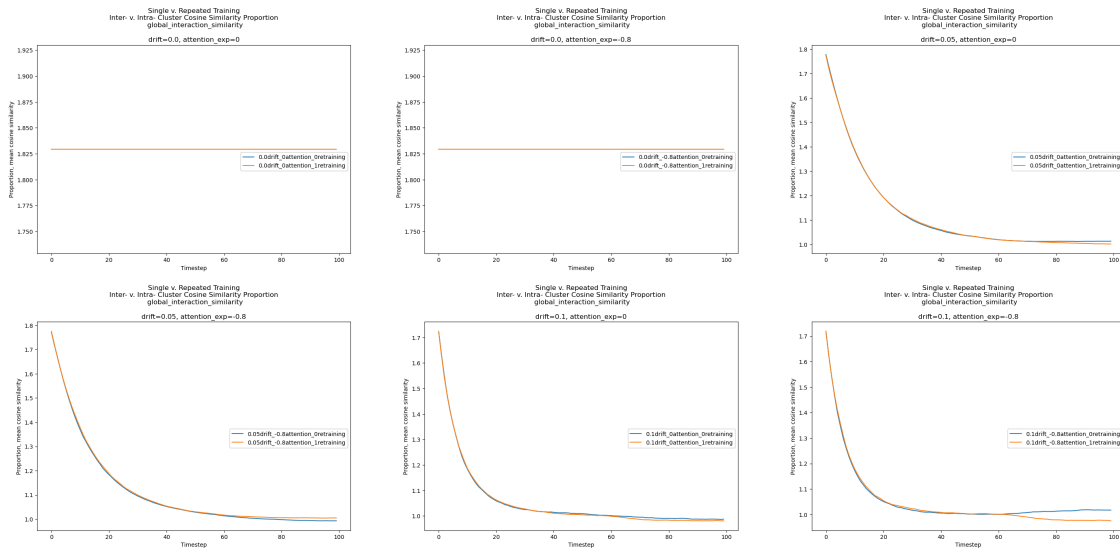
0.0drift_0attention_1retraining
0.0drift_-0.8attention_1retraining
0.05drift_0attention_1retraining
0.05drift_-0.8attention_1retraining
0.1drift_0attention_1retraining
0.1drift_-0.8attention_1retraining



```
[ ]: numerator = "mean_intra_cluster_cosine_sim"
denominator = "mean_inter_cluster_cosine_sim"

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, denominator, numerator,
    ↪model_key_pairs[i], id_to_readable, mult_sd=0)
    curr_ax.set_ylabel("Proportion, mean cosine similarity")
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Single v. Repeated Training\nInter- v. Intra- Cluster_\n↪Cosine Similarity Proportion\n↪{global_metric_key}\n↪ndrift={models[model_key_pairs[i][0]][0]},\n↪attention_exp={models[model_key_pairs[i][0]][1]}")
    # curr_ax.set_ylim(-0.1, 1.75)
    # curr_ax.set_xlim(40, 100)
```



4 Graphing Distance from centroid

4.0.1 Graphing mean_cluster_distance_from_centroid

```
[ ]: metric_key = "mean_cluster_distance_from_centroid"

single_training_keys = [key[0] for key in model_key_pairs]
repeated_training_keys = [key[1] for key in model_key_pairs]
```

```

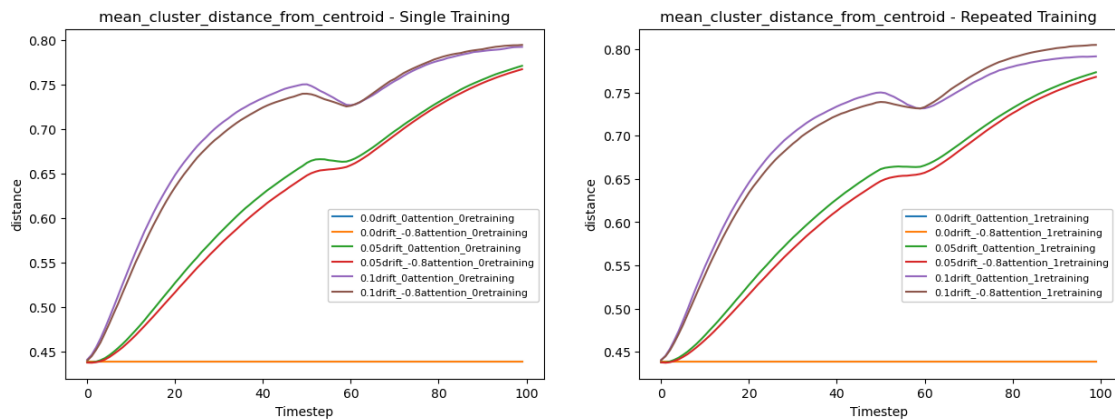
fig, axs = plt.subplots(1, 2, figsize=(15, 5))

graph_metrics_by_axis(axs[0], results, metric_key, single_training_keys,
    ↪id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right',
    ↪bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axs[1], results, metric_key, repeated_training_keys,
    ↪id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right',
    ↪bbox_to_anchor=(1, 0.5), fontsize="8",)

```

[]: <matplotlib.legend.Legend at 0x29a315610>



Graphing Intra-Cluster Distance relative to Global Distance

```

[ ]: """
    Single training
    """
    metric_keys = ['mean_cluster_distance_from_centroid',
        ↪'mean_global_distance_from_centroid']

    fig, axs = plt.subplots(2, 3, figsize=(30, 15))
    fig.tight_layout(pad=10.0)

    for i in range(len(model_key_pairs)):

```

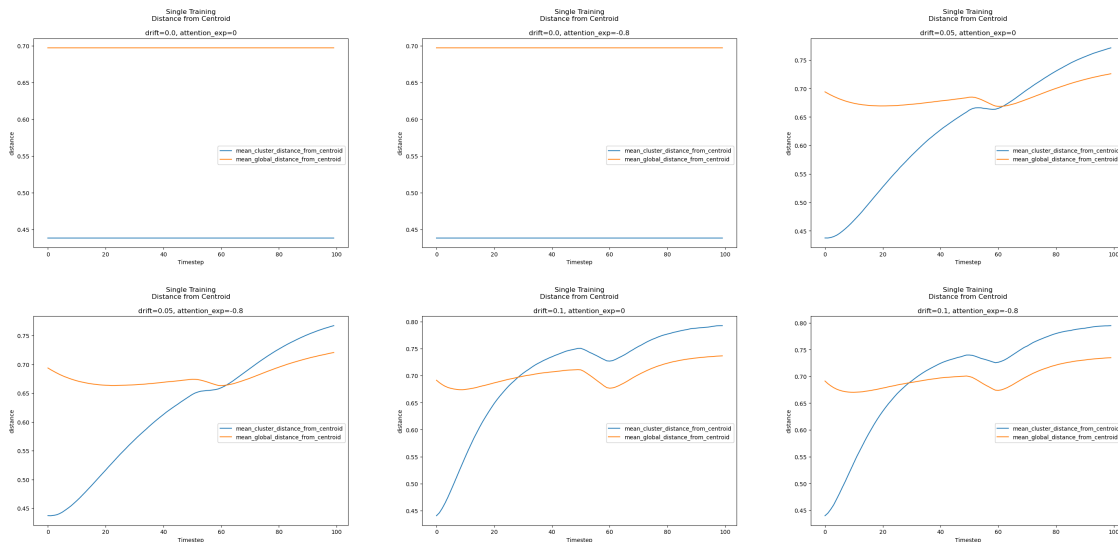
```

print(model_key_pairs[i][0])
curr_ax = axs[int(i >= 3), i%3]
graph_metrics_by_axis(curr_ax, results, metric_keys[0],
↳(model_key_pairs[i][0],), id_to_readable, mult_sd=0, graph_by="metric")
graph_metrics_by_axis(curr_ax, results, metric_keys[1],
↳(model_key_pairs[i][0],), id_to_readable, mult_sd=0, graph_by="metric")

# curr_ax.set_xlim(75, 100)
# if i > 3:
#     curr_ax.set_ylim(0.3, 0.6)
# elif i > 1:
#     curr_ax.set_ylim(0.4, 0.8)
# if i > 1:
#     curr_ax.set_ylim(0.3, 0.8)
curr_ax.set_ylabel(y_labels[metric_keys[0]])
curr_ax.set_xlabel("Timestep")
curr_ax.set_title(f"Single Training\nDistance from_
↳Centroid\n\ndrift={models[model_key_pairs[i][0]][0]},
↳attention_exp={models[model_key_pairs[i][0]][1]}")

```

0.0drift_0attention_0retraining
 0.0drift_-0.8attention_0retraining
 0.05drift_0attention_0retraining
 0.05drift_-0.8attention_0retraining
 0.1drift_0attention_0retraining
 0.1drift_-0.8attention_0retraining



```

[ ]: """
Repeated training

```

```

"""
metric_keys = ['mean_cluster_distance_from_centroid',
               ↪ 'mean_global_distance_from_centroid']

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    print(model_key_pairs[i][0])
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_keys[0],
    ↪ (model_key_pairs[i][1],), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[1],
    ↪ (model_key_pairs[i][1],), id_to_readable, mult_sd=0, graph_by="metric")

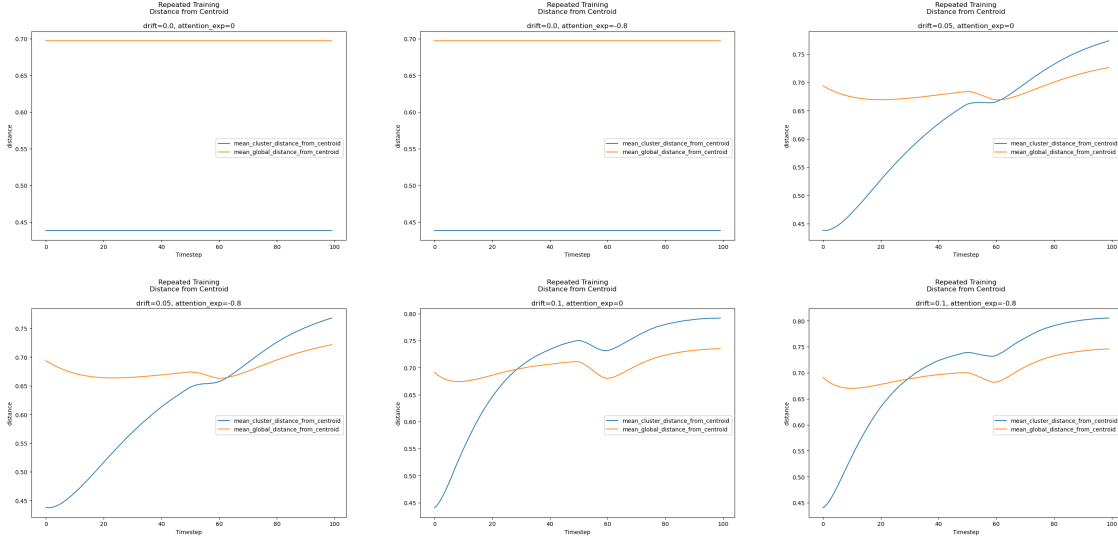
    # curr_ax.set_xlim(75, 100)
    # if i > 3:
    #     curr_ax.set_ylim(0.3, 0.6)
    # elif i > 1:
    #     curr_ax.set_ylim(0.4, 0.8)
    # if i > 1:
    #     curr_ax.set_ylim(0.3, 0.8)
    curr_ax.set_ylabel(y_labels[metric_keys[0]])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Repeated Training\nDistance from
    ↪ Centroid\n\ndrift={models[model_key_pairs[i][0]][0]},
    ↪ attention_exp={models[model_key_pairs[i][0]][1]}")

```

```

0.0drift_0attention_0retraining
0.0drift_-0.8attention_0retraining
0.05drift_0attention_0retraining
0.05drift_-0.8attention_0retraining
0.1drift_0attention_0retraining
0.1drift_-0.8attention_0retraining

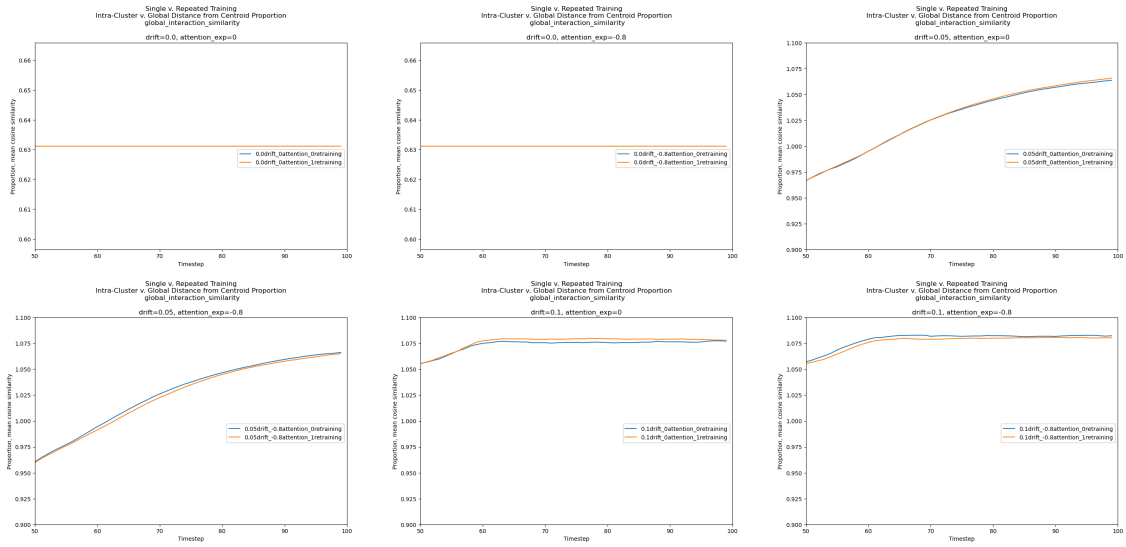
```



```
[ ]: numerator = "mean_cluster_distance_from_centroid"
denominator = "mean_global_distance_from_centroid"

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, denominator, numerator,
    ↪model_key_pairs[i], id_to_readable, mult_sd=0)
    curr_ax.set_xlim(50, 100)
    if i > 1:
        curr_ax.set_ylim(0.9, 1.1)
        curr_ax.set_ylabel("Proportion, mean cosine similarity")
        curr_ax.set_xlabel("Timestep")
        curr_ax.set_title(f"Single v. Repeated Training\nIntra-Cluster v. Global_
        ↪Distance from Centroid Proportion\n_
        ↪{global_metric_key}\n\ndrift={models[model_key_pairs[i][0]][0]},_
        ↪attention_exp={models[model_key_pairs[i][0]][1]}")
        # curr_ax.set_ylim(-0.1, 1.75)
        # curr_ax.set_xlim(40, 100)
```



[]: