

param_exp_results-with_clustering_metrics-50train50run

June 22, 2023

```
[ ]: import sys
sys.path.insert(1, '/Users/madisonthantu/Desktop/DREAM/t-recs')

import trecs

import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
from scipy.ndimage import gaussian_filter1d
from collections import defaultdict
from chaney_utils import (
    load_sim_results,
    graph_relative_to_ideal,
    merge_results,
    graph_metrics,
    graph_metrics_by_axis,
    graph_relative_to_global_by_axis,
    transform_relative_to_global,
    graph_histogram_metric_by_axis
)
import warnings
warnings.simplefilter("ignore")

import itertools
```

```
/Users/madisonthantu/Desktop/T-RECS-RS-
research/prelim_experiments/param_experiments/chaney_utils.py:61: SyntaxWarning:
"is" with a literal. Did you mean "=="?
    if model_key is "ideal" and not absolute_measure:
```

```
[ ]: DEBUG = True # don't save figures
```

```
[ ]: # results = merge_results(["param_exp_results/repeated_training", ↴
    "param_exp_results/single_training"])

results_path = "param_exp_results/with_clustering_metrics/10train90run"
```

```

results = merge_results([results_path])
print(results.keys())

model_keys = ['0.0drift_0attention_0retraining', '0.
↳0drift_0attention_1retraining', '0.0drift_-0.8attention_0retraining', '0.
↳0drift_-0.8attention_1retraining', '0.05drift_0attention_0retraining', '0.
↳05drift_0attention_1retraining', '0.05drift_-0.8attention_0retraining', '0.
↳05drift_-0.8attention_1retraining', '0.1drift_0attention_0retraining', '0.
↳1drift_0attention_1retraining', '0.1drift_-0.8attention_0retraining', '0.
↳1drift_-0.8attention_1retraining']

print(model_keys)

id_to_readable = dict(zip(model_keys, model_keys))
print(id_to_readable)

```

```

dict_keys(['mse', 'interaction_spread', 'global_interaction_similarity',
'inter_cluster_interaction_similarity', 'intra_cluster_interaction_similarity',
'mean_global_cosine_sim', 'mean_intra_cluster_cosine_sim',
'mean_inter_cluster_cosine_sim', 'mean_cosine_sim_per_cluster',
'mean_cluster_distance_from_centroid', 'mean_global_distance_from_centroid',
'mean_distance_from_centroid_per_cluster'])

['0.0drift_0attention_0retraining', '0.0drift_0attention_1retraining',
'0.0drift_-0.8attention_0retraining', '0.0drift_-0.8attention_1retraining',
'0.05drift_0attention_0retraining', '0.05drift_0attention_1retraining',
'0.05drift_-0.8attention_0retraining', '0.05drift_-0.8attention_1retraining',
'0.1drift_0attention_0retraining', '0.1drift_0attention_1retraining',
'0.1drift_-0.8attention_0retraining', '0.1drift_-0.8attention_1retraining']

{'0.0drift_0attention_0retraining': '0.0drift_0attention_0retraining',
'0.0drift_0attention_1retraining': '0.0drift_0attention_1retraining',
'0.0drift_-0.8attention_0retraining': '0.0drift_-0.8attention_0retraining',
'0.0drift_-0.8attention_1retraining': '0.0drift_-0.8attention_1retraining',
'0.05drift_0attention_0retraining': '0.05drift_0attention_0retraining',
'0.05drift_0attention_1retraining': '0.05drift_0attention_1retraining',
'0.05drift_-0.8attention_0retraining': '0.05drift_-0.8attention_0retraining',
'0.05drift_-0.8attention_1retraining': '0.05drift_-0.8attention_1retraining',
'0.1drift_0attention_0retraining': '0.1drift_0attention_0retraining',
'0.1drift_0attention_1retraining': '0.1drift_0attention_1retraining',
'0.1drift_-0.8attention_0retraining': '0.1drift_-0.8attention_0retraining',
'0.1drift_-0.8attention_1retraining': '0.1drift_-0.8attention_1retraining'}

```

```

[ ]: hyper_params = {"drift": [0.0, 0.05, 0.1], "attention_exp": [0, -0.8], ↳
↳"repeated_training": [0, 1]}

models = dict([(f"{p[0]}drift_{p[1]}attention_{p[2]}retraining", p) for p in ↳
↳itertools.product(*hyper_params.values())])

results_df = pd.DataFrame(columns=["drift", "attention_exp", ↳
↳"repeated_training"] + list(results.keys()))

```

```

results_df

for params in model_keys:
    df = pd.DataFrame(columns=["drift", "attention_exp", "repeated_training"] +_
list(results.keys()))
    for metric in results:
        metric_results = [vals for sim_trial in results[metric][params] for_
vals in sim_trial]
        df[metric] = metric_results
    df["drift"] = models[params][0]
    df["attention_exp"] = models[params][1]
    df["repeated_training"] = models[params][2]
    results_df = pd.concat([results_df, df])

```

```
[ ]: print(results_df.shape)
results_df.head()
# len(models)
```

(3600, 15)

```
[ ]:   drift attention_exp repeated_training      mse interaction_spread \
0     0.0          0            0  0.090065           -938.5
1     0.0          0            0  0.090065            0.0
2     0.0          0            0  0.090065           -1.0
3     0.0          0            0  0.090065            1.0
4     0.0          0            0  0.090065            0.0

      global_interaction_similarity  inter_cluster_interaction_similarity \
0                  0.001601           0.001386
1                  0.002274           0.002032
2                  0.003058           0.002794
3                  0.003888           0.003562
4                  0.004737           0.004383

      intra_cluster_interaction_similarity  mean_global_cosine_sim \
0                  0.001954           0.244844
1                  0.002672           0.244844
2                  0.003493           0.244844
3                  0.004424           0.244844
4                  0.005318           0.244844

      mean_intra_cluster_cosine_sim  mean_inter_cluster_cosine_sim \
0                  0.318086           0.200335
1                  0.318086           0.200335
2                  0.318086           0.200335
3                  0.318086           0.200335
4                  0.318086           0.200335
```

```

                mean_cosine_sim_per_cluster \
0 [0.5116548734034113, 0.9312028843954681, 0.727...
1 [0.5116548734034113, 0.9312028843954681, 0.727...
2 [0.5116548734034113, 0.9312028843954681, 0.727...
3 [0.5116548734034113, 0.9312028843954681, 0.727...
4 [0.5116548734034113, 0.9312028843954681, 0.727...

mean_cluster_distance_from_centroid  mean_global_distance_from_centroid \
0                               0.43184                      0.687343
1                               0.43184                      0.687343
2                               0.43184                      0.687343
3                               0.43184                      0.687343
4                               0.43184                      0.687343

mean_distance_from_centroid_per_cluster
0 [0.46841372717852575, 0.777901889986511, 0.688...
1 [0.46841372717852575, 0.777901889986511, 0.688...
2 [0.46841372717852575, 0.777901889986511, 0.688...
3 [0.46841372717852575, 0.777901889986511, 0.688...
4 [0.46841372717852575, 0.777901889986511, 0.688...

[ ]: y_labels = dict([
    ("mse", "mse"),
    ("interaction_spread", "interaction spread"),
    ("global_interaction_similarity", "average jacard similarity"),
    ("inter_cluster_interaction_similarity", "average jacard similarity"),
    ("intra_cluster_interaction_similarity", "average jacard similarity"),
    ("mean_global_cosine_sim", "average cosine sim"),
    ("mean_intra_cluster_cosine_sim", "average cosine sim"),
    ("mean_inter_cluster_cosine_sim", "average cosine sim"),
    ("mean_cosine_sim_per_cluster", "average cosine sim"),
    ("mean_cluster_distance_from_centroid", "distance"),
    ("mean_global_distance_from_centroid", "distance"),
    ("mean_distance_from_centroid_per_cluster", "distance"),
])
model_key_pairs = [(model_keys[i], model_keys[i+1]) for i in range(0, len(model_keys), 2)]

single_training_keys = [key[0] for key in model_key_pairs]
repeated_training_keys = [key[1] for key in model_key_pairs]
```

1 Graphing MSE

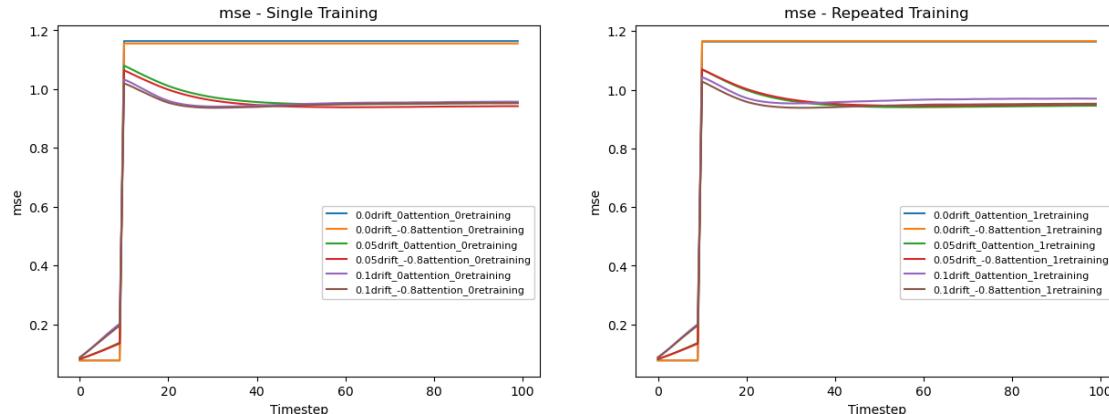
```
[ ]: metric_key = "mse"

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

graph_metrics_by_axis(axs[0], results, metric_key, single_training_keys, ↴
    id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right', ↴
    bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axs[1], results, metric_key, repeated_training_keys, ↴
    id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right', ↴
    bbox_to_anchor=(1, 0.5), fontsize="8",)
```

```
[ ]: <matplotlib.legend.Legend at 0x16a09abb0>
```



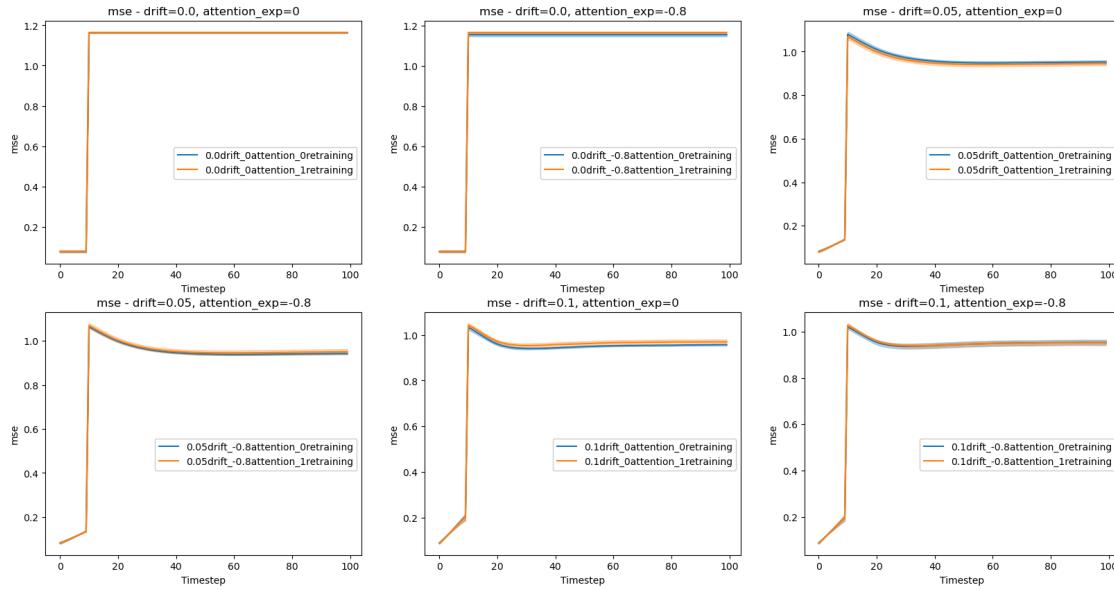
```
[ ]: metric_key = "mse"

fig, axs = plt.subplots(2, 3, figsize=(20, 10))

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_key, model_key_pairs[i], ↴
        id_to_readable, mult_sd=0.5)
```

```

curr_ax.set_ylabel(y_labels[metric_key])
curr_ax.set_xlabel("Timestep")
curr_ax.set_title(f"{metric_key} - "
↳drift={models[model_key_pairs[i][0]][0]}, ↳
↳attention_exp={models[model_key_pairs[i][0]][1]})
```



```

[ ]: # single_training_keys = ['0.Odrift_Oattention_Oretraining', '0.Odrift_-0.
↳8attention_Oretraining', '0.05drift_Oattention_Oretraining', '0.05drift_-0.
↳8attention_Oretraining', '0.1drift_Oattention_Oretraining', '0.1drift_-0.
↳8attention_Oretraining']
# print(single_training_keys)

# metric_key = "mse"
# # graph_relative_to_ideal(results, "mse", single_training_keys, ↳
# ↳id_to_readable, absolute_measure=False, mult_sd=1.0)
# graph_metrics(results, metric_key, single_training_keys, id_to_readable, ↳
# ↳mult_sd=0.5)
# plt.ylabel(y_labels[metric_key])
# plt.xlabel("Timestep")
# plt.legend(facecolor='white', framealpha=1, loc='upper right', ↳
# ↳bbox_to_anchor=(1.7, 1.0))
# plt.ylim(-0.1, 1.75)
# plt.xlim(0, 60)
# plt.title(f"Single Training - {metric_key}")
# if not DEBUG:
#     plt.savefig("figures/repeated_training_sim_pair.pdf", bbox_inches = ↳
# ↳"tight")
```

```

# repeated_training_keys = list(set(model_keys) - set(single_training_keys))
# print(repeated_training_keys)

# metric_key = "mse"
# # graph_relative_to_ideal(results, "mse", single_training_keys, id_to_readable, absolute_measure=False, mult_sd=1.0)
# graph_metrics(results, metric_key, repeated_training_keys, id_to_readable, mult_sd=0.5)
# plt.ylabel("Change in Jaccard Index")
# plt.xlabel("Timestep")
# plt.legend(facecolor='white', framealpha=1, loc='upper right', bbox_to_anchor=(1.7, 1.0))
# plt.ylim(-0.1, 1.75)
# plt.xlim(0, 60)
# plt.title(f'Repeated Training - {metric_key}')
# if not DEBUG:
#     plt.savefig("figures/repeated_training_sim_pair.pdf", bbox_inches = "tight")

```

2 Graphing Interaction Similarity

2.0.1 Graphing Global Interaction Similarity

```

[ ]: metric_key = "global_interaction_similarity"

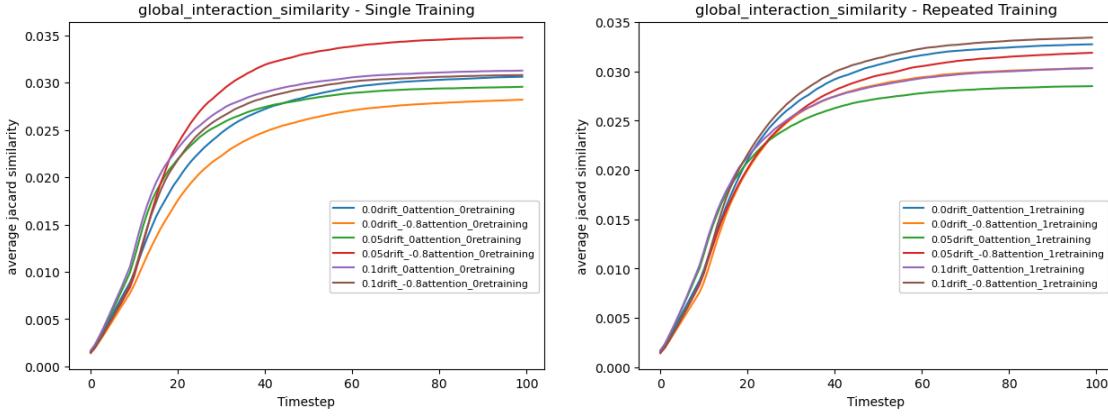
fig, axs = plt.subplots(1, 2, figsize=(15, 5))

graph_metrics_by_axis(axs[0], results, metric_key, single_training_keys, id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right', bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axs[1], results, metric_key, repeated_training_keys, id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right', bbox_to_anchor=(1, 0.5), fontsize="8",)

```

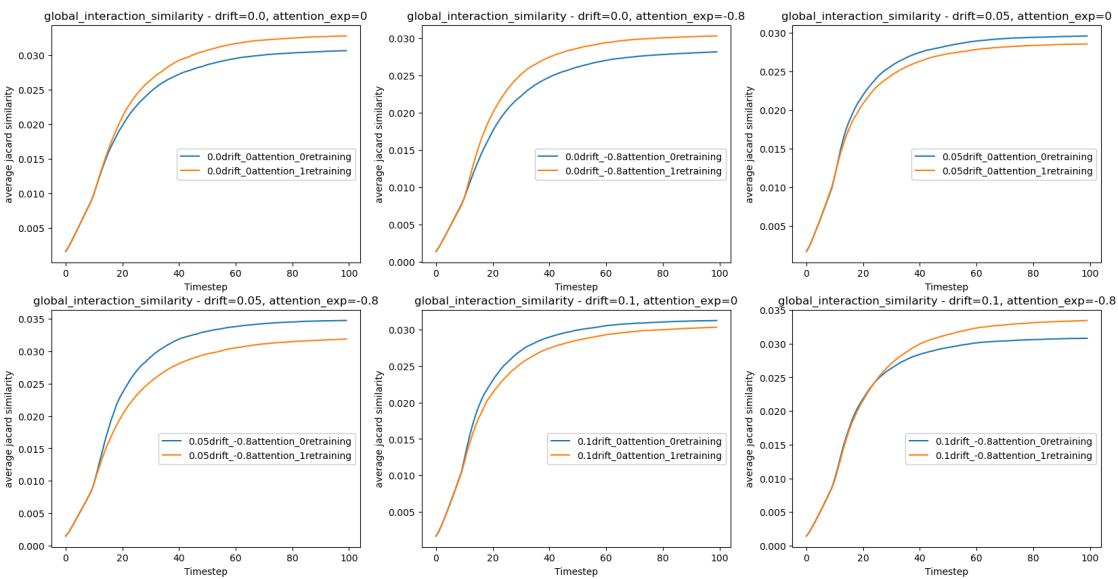
```
[ ]: <matplotlib.legend.Legend at 0x16b281d00>
```



```
[ ]: metric_key = "global_interaction_similarity"

fig, axs = plt.subplots(2, 3, figsize=(20, 10))

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_key, model_key_pairs[i], id_to_readable, mult_sd=0)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"{metric_key} - drift={models[model_key_pairs[i][0]][0]}, attention_exp={models[model_key_pairs[i][0]][1]}")
```



2.0.2 Graphing Inter -Cluster Interaction Similarity

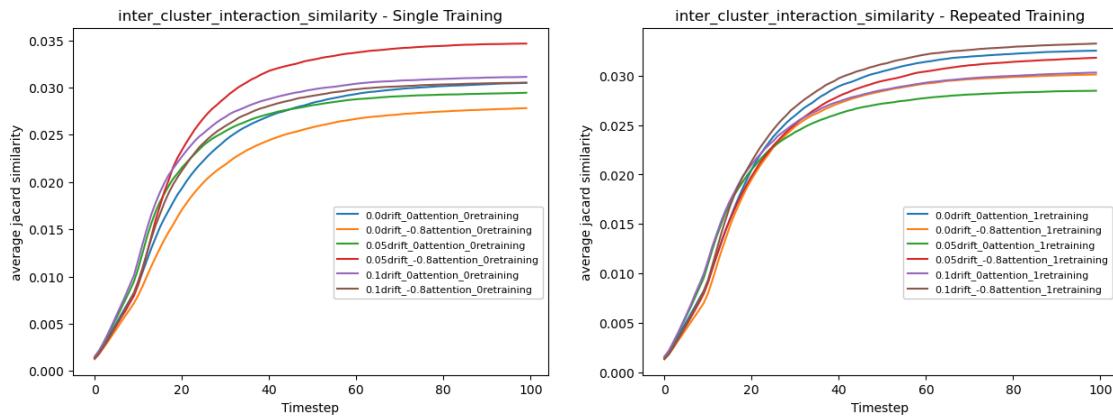
```
[ ]: metric_key = "inter_cluster_interaction_similarity"

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

graph_metrics_by_axis(axs[0], results, metric_key, single_training_keys, ▾
    ↵id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right', ▾
    ↵bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axs[1], results, metric_key, repeated_training_keys, ▾
    ↵id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right', ▾
    ↵bbox_to_anchor=(1, 0.5), fontsize="8",)
```

[]: <matplotlib.legend.Legend at 0x16a862a90>



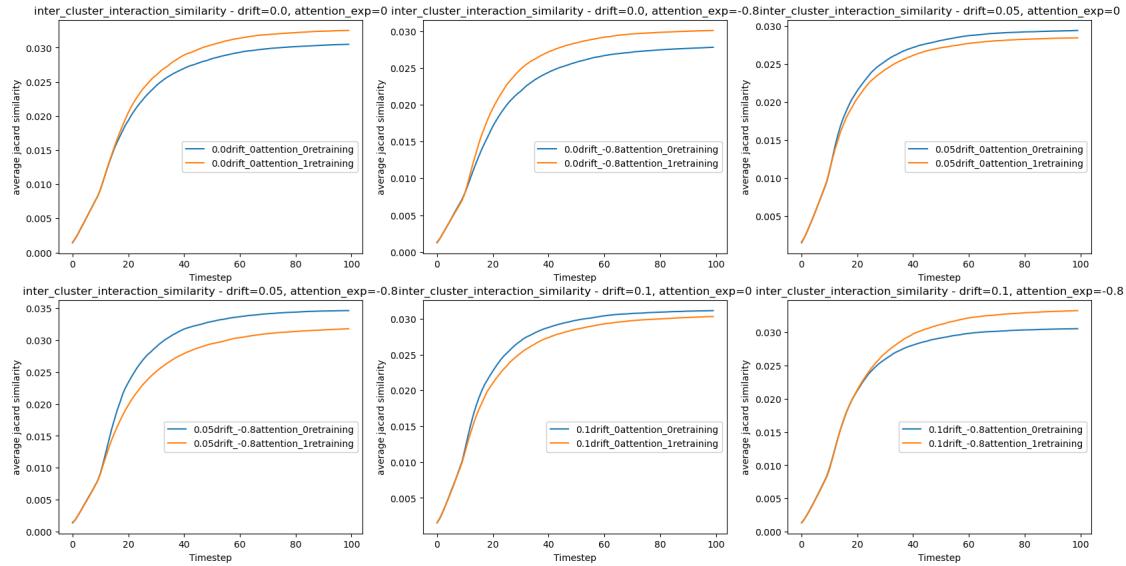
```
[ ]: metric_key = "inter_cluster_interaction_similarity"

fig, axs = plt.subplots(2, 3, figsize=(20, 10))

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_key, model_key_pairs[i], ▾
        ↵id_to_readable, mult_sd=0)
```

```

curr_ax.set_ylabel(y_labels[metric_key])
curr_ax.set_xlabel("Timestep")
curr_ax.set_title(f"{metric_key} - "
↳drift={models[model_key_pairs[i][0]][0]}, ↳
↳attention_exp={models[model_key_pairs[i][0]][1]}")
```



2.0.3 Graphing Intra -Cluster Interaction Similarity

```

[ ]: metric_key = "intra_cluster_interaction_similarity"

single_training_keys = [key[0] for key in model_key_pairs]
repeated_training_keys = [key[1] for key in model_key_pairs]

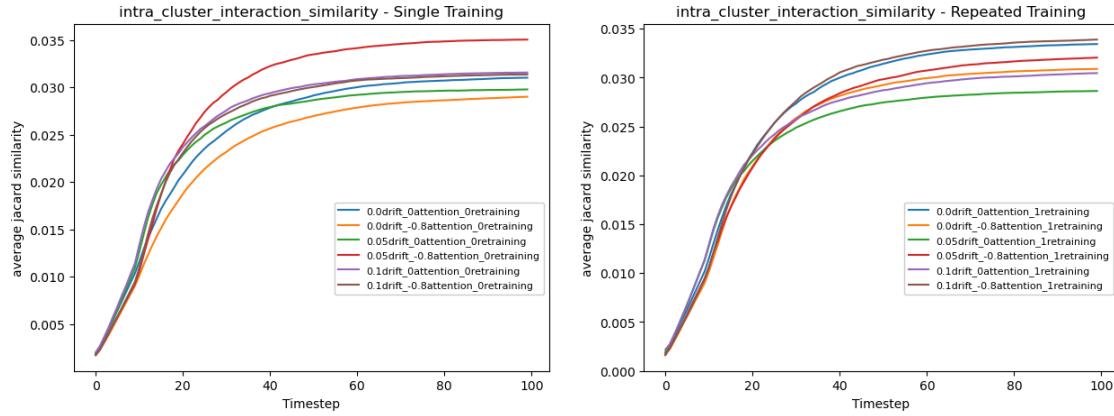
fig, axs = plt.subplots(1, 2, figsize=(15, 5))

graph_metrics_by_axis(axs[0], results, metric_key, single_training_keys, ↳
↳id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right', ↳
↳bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axs[1], results, metric_key, repeated_training_keys, ↳
↳id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
```

```
axs[1].legend(facecolor='white', framealpha=1, loc='upper right',  
bbox_to_anchor=(1, 0.5), fontsize="8",)
```

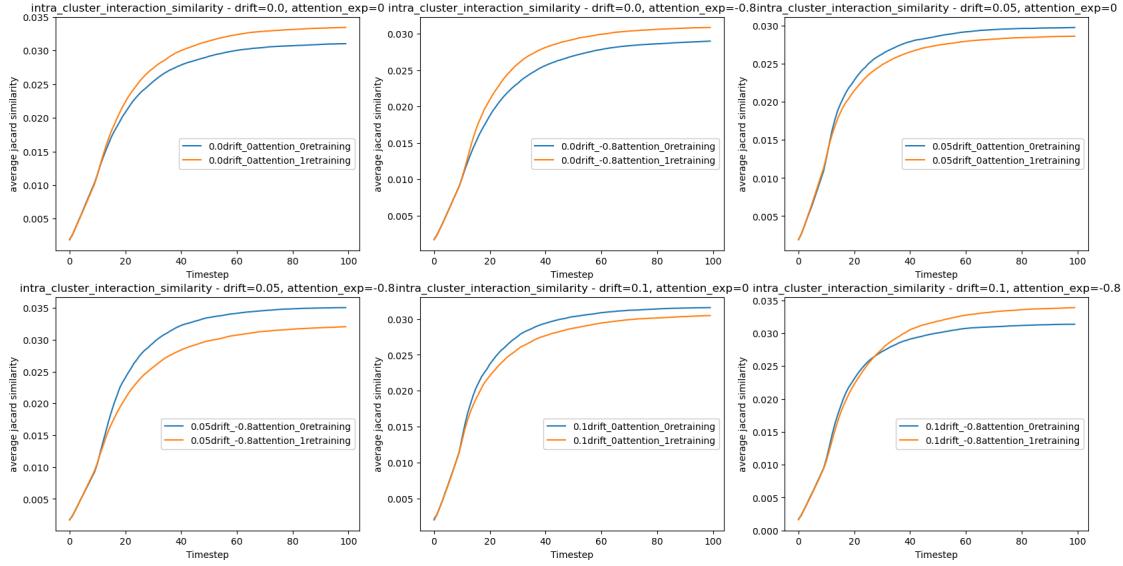
[]: <matplotlib.legend.Legend at 0x16b6effd0>



[]: metric_key = "intra_cluster_interaction_similarity"

```
fig, axs = plt.subplots(2, 3, figsize=(20, 10))

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_key, model_key_pairs[i],  
    ↪id_to_readable, mult_sd=0)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"{metric_key} -  
    ↪drift={models[model_key_pairs[i][0]][0]},  
    ↪attention_exp={models[model_key_pairs[i][0]][1]}")
```



2.0.4 Graphing Intra -Cluster Interaction Similarity and Inter -Cluster Interaction Similarity relative to Global Interaction Simarilty

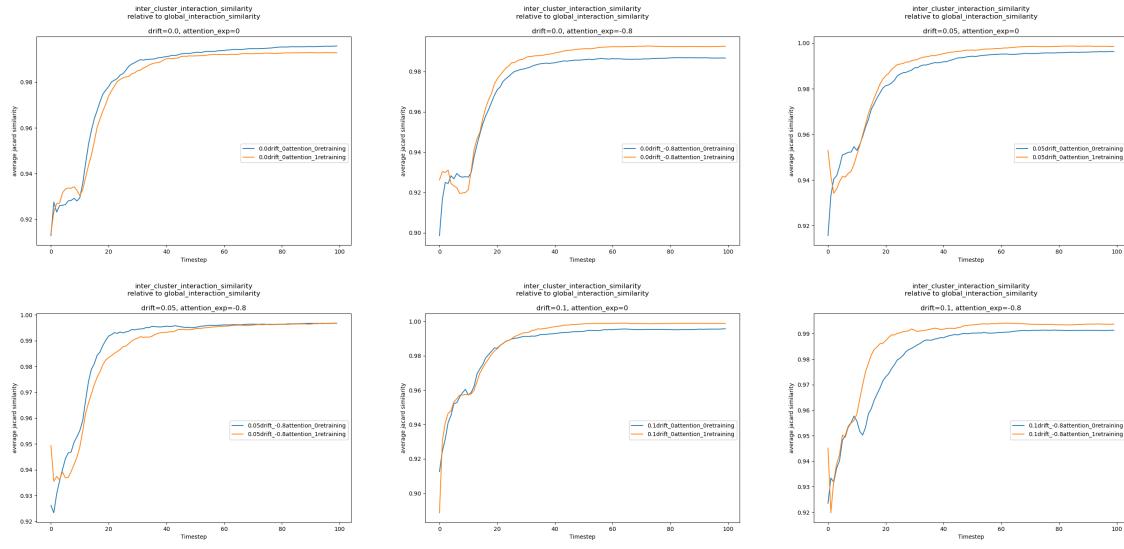
```
[ ]: global_metric_key = "global_interaction_similarity"
```

Graphing Inter -Cluster Interaction Similarity relative to Global Interaction Simarilty Single training vs. Retraining

```
[ ]: metric_key = "inter_cluster_interaction_similarity"

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
    ↪metric_key, model_key_pairs[i], id_to_readable, mult_sd=0)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"{metric_key}\nrelative to\n{global_metric_key}\nndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")
```



```
[ ]: metric_key = "inter_cluster_interaction_similarity"

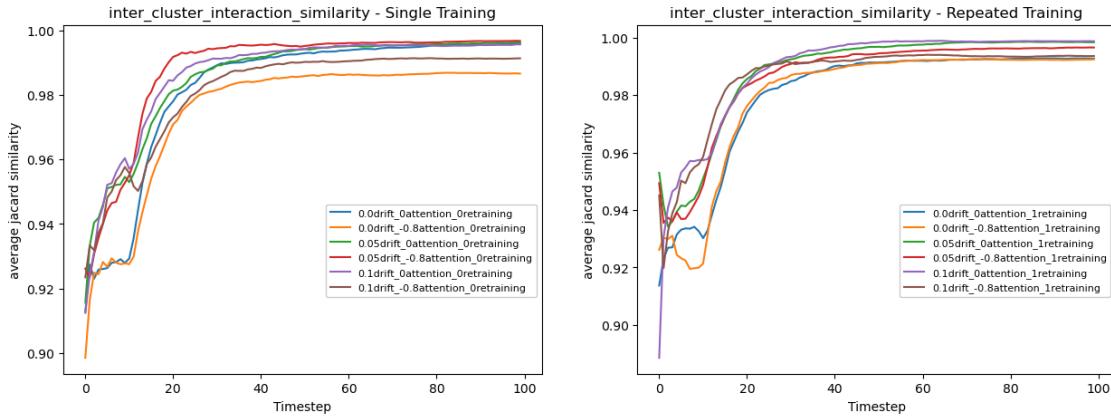
single_training_keys = [key[0] for key in model_key_pairs]
repeated_training_keys = [key[1] for key in model_key_pairs]

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

# graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
#                                  metric_key, model_key_pairs[i], id_to_readable, mult_sd=0)
graph_relative_to_global_by_axis(axs[0], results, global_metric_key,
                                 metric_key, single_training_keys, id_to_readable, mult_sd=0)
graph_relative_to_global_by_axis(axs[1], results, global_metric_key,
                                 metric_key, repeated_training_keys, id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right',
              bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_relative_to_global_by_axis(axs[1], results, global_metric_key,
                                 metric_key, repeated_training_keys, id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right',
              bbox_to_anchor=(1, 0.5), fontsize="8",)
```

```
[ ]: <matplotlib.legend.Legend at 0x16b153dc0>
```



```
[ ]: # inter_cluster_interaction_similarity = results['inter_cluster_interaction_similarity']['0.1drift_-0.8attention_1retraining'].mean(axis=0)
# global_interaction_similarity = results['global_interaction_similarity']['0.1drift_-0.8attention_1retraining'].mean(axis=0)

# diff = np.divide(inter_cluster_interaction_similarity, global_interaction_similarity)

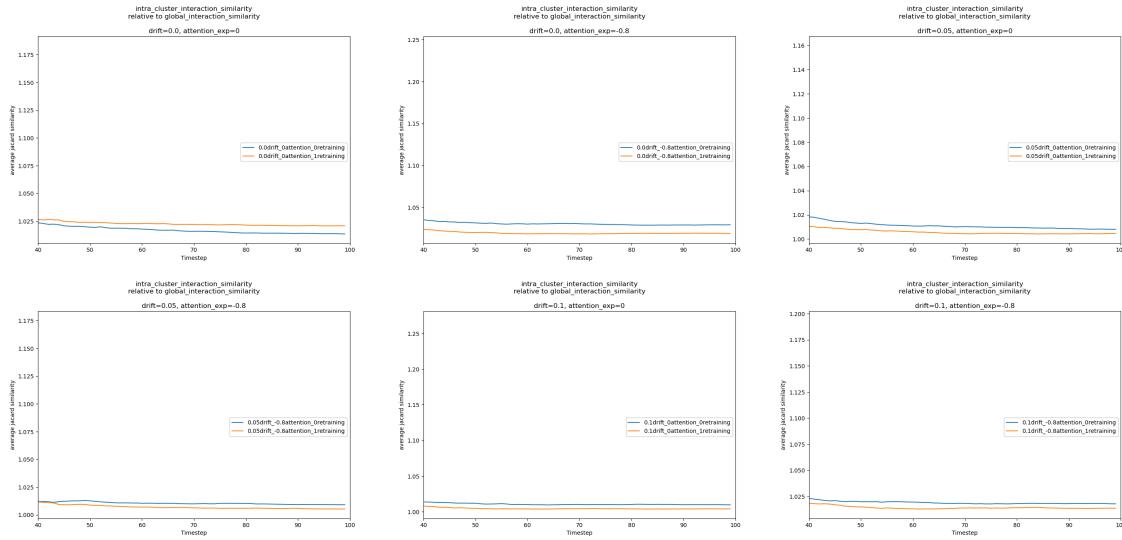
# plt.plot(diff)
```

Graphing Intra -Cluster Interaction Similarity relative to Global Interaction Simarilty Single training vs. Retraining

```
[ ]: metric_key = "intra_cluster_interaction_similarity"

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, global_metric_key, metric_key, model_key_pairs[i], id_to_readable, mult_sd=0)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"{metric_key}\nrelative to {global_metric_key}\ndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")
    # curr_ax.set_ylim(-0.1, 1.75)
    curr_ax.set_xlim(40, 100)
```



```
[ ]: metric_key = "intra_cluster_interaction_similarity"

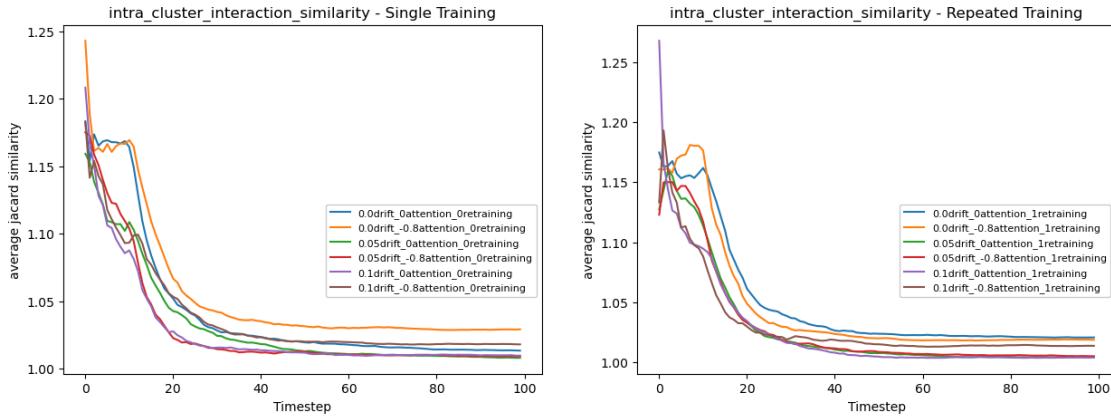
single_training_keys = [key[0] for key in model_key_pairs]
repeated_training_keys = [key[1] for key in model_key_pairs]

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

# graph_relative_to_global_by_axis(curr_ax, results, global_metric_key, ↴
# metric_key, model_key_pairs[i], id_to_readable, mult_sd=0)
graph_relative_to_global_by_axis(axs[0], results, global_metric_key, ↴
metric_key, single_training_keys, id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right', ↴
bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_relative_to_global_by_axis(axs[1], results, global_metric_key, ↴
metric_key, repeated_training_keys, id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right', ↴
bbox_to_anchor=(1, 0.5), fontsize="8",)

[ ]: <matplotlib.legend.Legend at 0x16b37eee0>
```



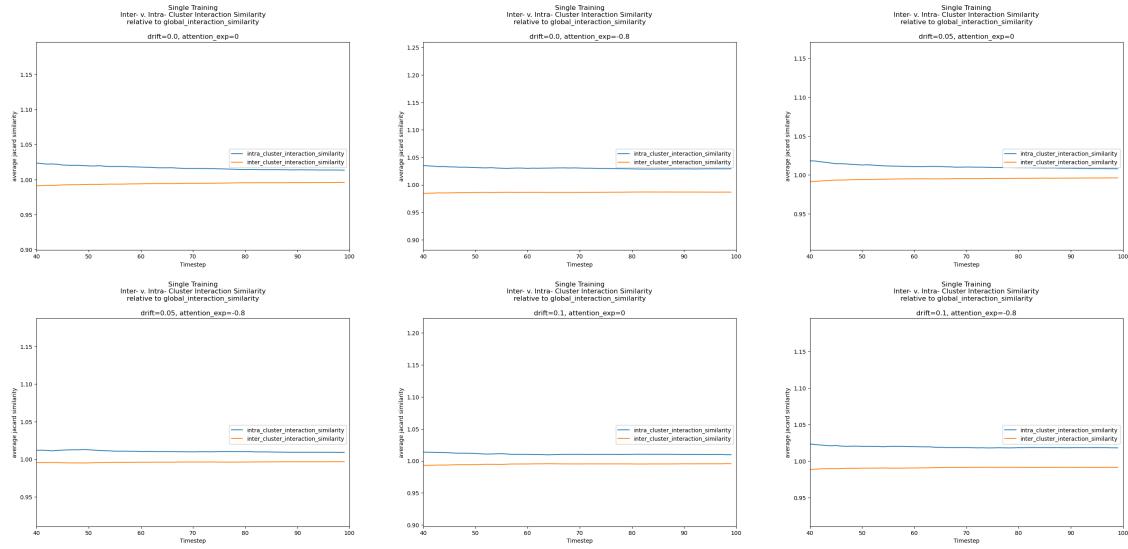
Graphing Cluster Interaction Similarity relative to Global Interaction Simarility
Inter- vs. Intra-

```
[ ]: """
Single training
"""

metric_keys = ["intra_cluster_interaction_similarity", "inter_cluster_interaction_similarity"]

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
                                     metric_keys[0], (model_key_pairs[i][0]), id_to_readable, mult_sd=0,
                                     graph_by="metric")
    graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
                                     metric_keys[1], (model_key_pairs[i][0]), id_to_readable, mult_sd=0,
                                     graph_by="metric")
    curr_ax.set_ylabel(y_labels[metric_keys[0]])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Single Training\nInter- v. Intra- Cluster Interaction\nSimilarity\nrelative to\n{global_metric_key}\nndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")
    # curr_ax.set_ylim(-0.1, 1.75)
    curr_ax.set_xlim(40, 100)
```

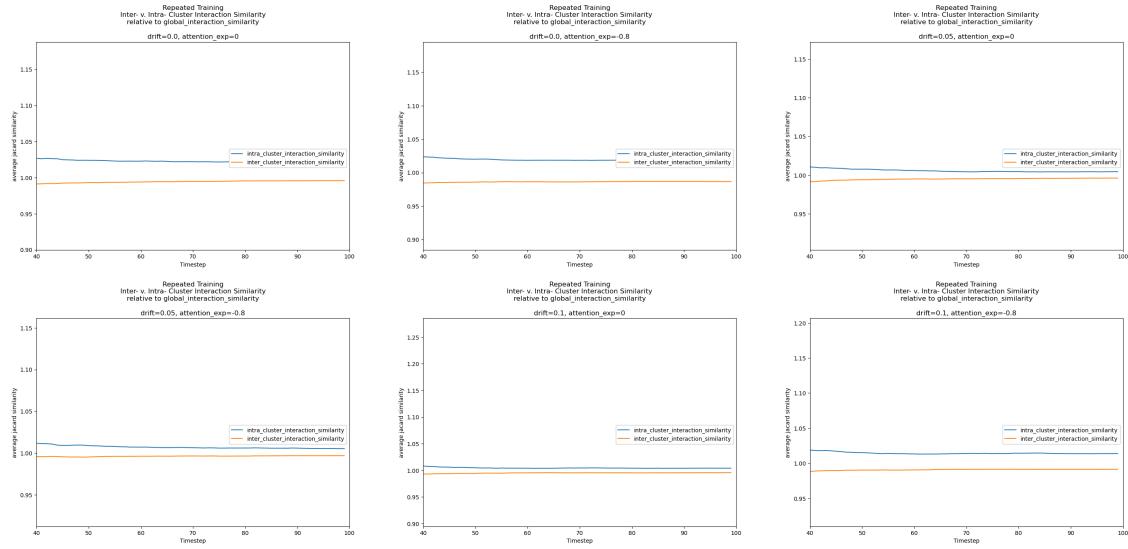


```
[ ]: """
Repeated training
"""

metric_keys = ["intra_cluster_interaction_similarity", "inter_cluster_interaction_similarity"]

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
                                    metric_keys[0], (model_key_pairs[i][1],), id_to_readable, mult_sd=0,
                                    graph_by="metric")
    graph_relative_to_global_by_axis(curr_ax, results, global_metric_key,
                                    metric_keys[1], (model_key_pairs[i][0],), id_to_readable, mult_sd=0,
                                    graph_by="metric")
    curr_ax.set_ylabel(y_labels[metric_keys[0]])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Repeated Training\nInter- v. Intra- Cluster Interaction\nSimilarity\nrelative to\n{global_metric_key}\n\ndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")
    # curr_ax.set_ylim(-0.1, 1.75)
    curr_ax.set_xlim(40, 100)
```

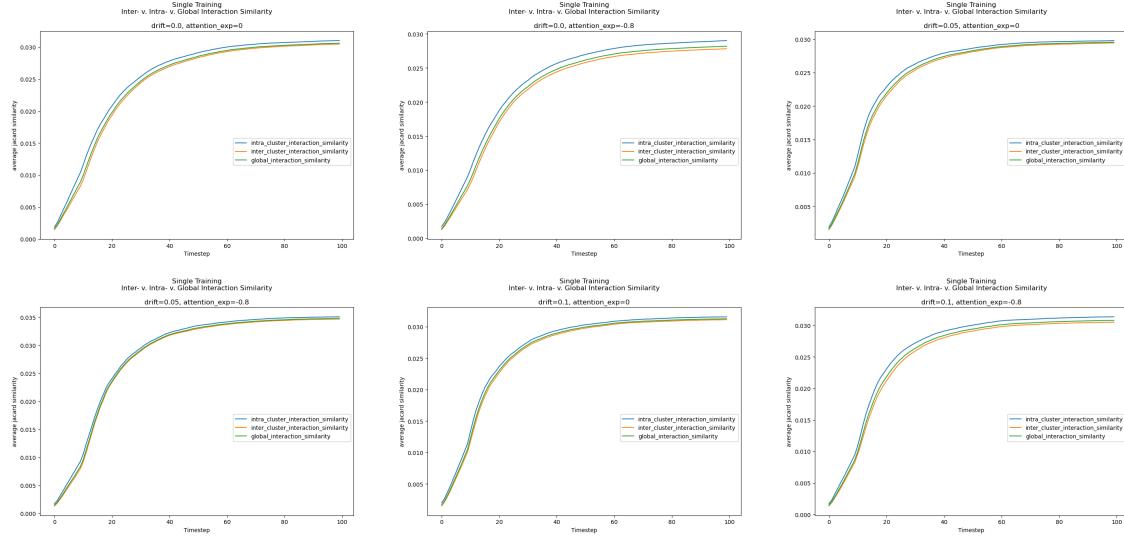


```
[ ]: """
Single training
"""

metric_keys = ["intra_cluster_interaction_similarity", "inter_cluster_interaction_similarity", "global_interaction_similarity"]

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_keys[0], model_key_pairs[i][0], id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[1], model_key_pairs[i][0], id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[2], model_key_pairs[i][0], id_to_readable, mult_sd=0, graph_by="metric")
    curr_ax.set_ylabel(y_labels[metric_keys[0]])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Single Training\nInter- v. Intra- v. Global Interaction\nSimilarity\nndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")
```



[]:

```
"""
Repeated training
"""

metric_keys = ["intra_cluster_interaction_similarity",
               "inter_cluster_interaction_similarity", "global_interaction_similarity"]

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    print(model_key_pairs[i][1])
    graph_metrics_by_axis(curr_ax, results, metric_keys[0],
                          (model_key_pairs[i][1],), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[1],
                          (model_key_pairs[i][1],), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[2],
                          (model_key_pairs[i][1],), id_to_readable, mult_sd=0, graph_by="metric")
    curr_ax.set_ylabel(y_labels[metric_keys[0]])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Repeated Training\nInter- v. Intra- v. Global\nInteraction Similarity\nndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")
```

Intra-Cluster:Inter-Cluster Interaction Similarity, Single training v. Repeated training

[]:

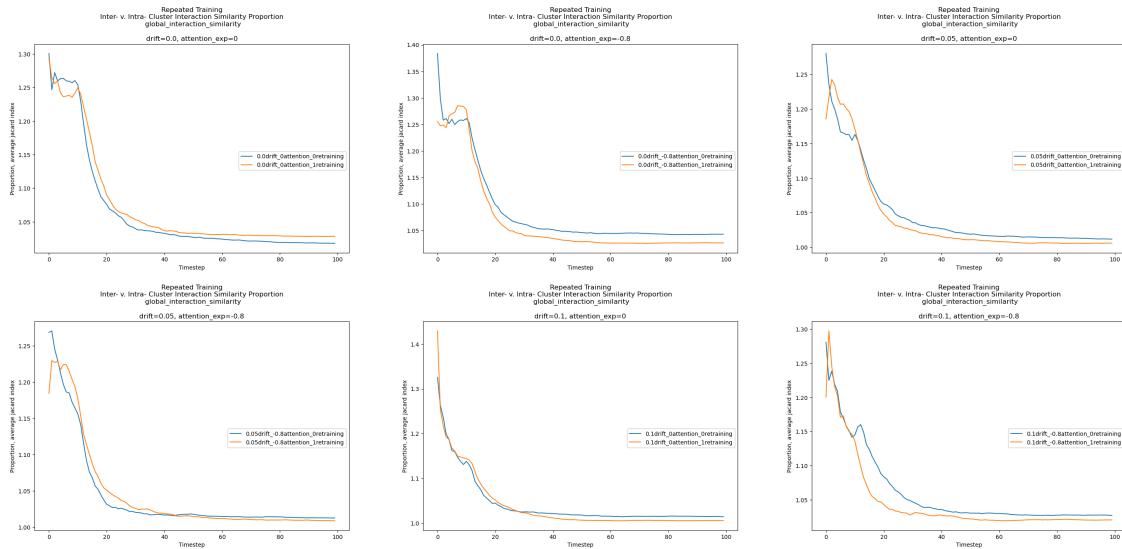
```
numerator = "intra_cluster_interaction_similarity"
denominator = "inter_cluster_interaction_similarity"
```

```

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, denominator, numerator, model_key_pairs[i], id_to_readable, mult_sd=0)
    curr_ax.set_ylabel("Proportion, average jacard index")
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Repeated Training\nInter- v. Intra- Cluster Interaction Similarity Proportion\n{global_metric_key}\n\nndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")
    # curr_ax.set_ylim(-0.1, 1.75)
    # curr_ax.set_xlim(40, 100)

```



3 Graphing Cosine Similarity

3.0.1 Graphing Intra -Cluster Cosine Similarity

```

[ ]: metric_key = "mean_intra_cluster_cosine_sim"

single_training_keys = [key[0] for key in model_key_pairs]
repeated_training_keys = [key[1] for key in model_key_pairs]

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

```

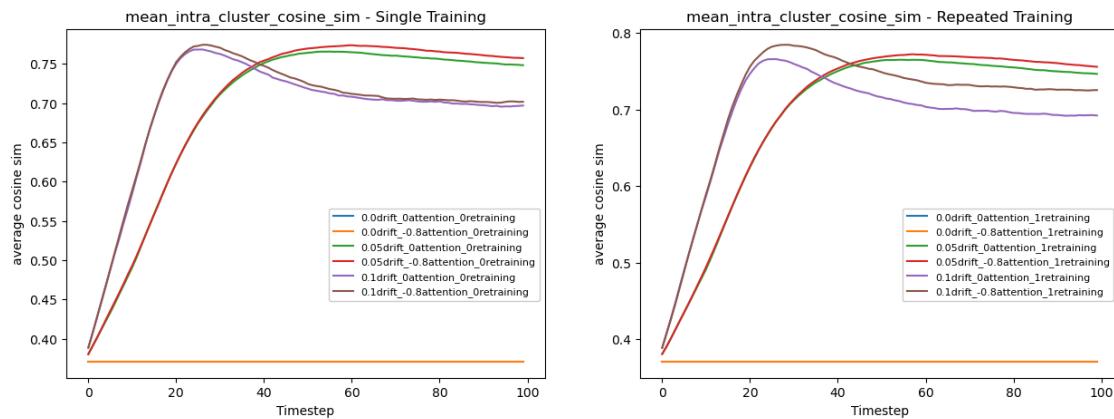
```

graph_metrics_by_axis(axes[0], results, metric_key, single_training_keys, ↴
    id_to_readable, mult_sd=0)
axes[0].set_ylabel(y_labels[metric_key])
axes[0].set_xlabel("Timestep")
axes[0].set_title(f"{metric_key} - Single Training")
axes[0].legend(facecolor='white', framealpha=1, loc='upper right', ↴
    bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axes[1], results, metric_key, repeated_training_keys, ↴
    id_to_readable, mult_sd=0)
axes[1].set_ylabel(y_labels[metric_key])
axes[1].set_xlabel("Timestep")
axes[1].set_title(f"{metric_key} - Repeated Training")
axes[1].legend(facecolor='white', framealpha=1, loc='upper right', ↴
    bbox_to_anchor=(1, 0.5), fontsize="8",)

```

[]: <matplotlib.legend.Legend at 0x16b55bb80>



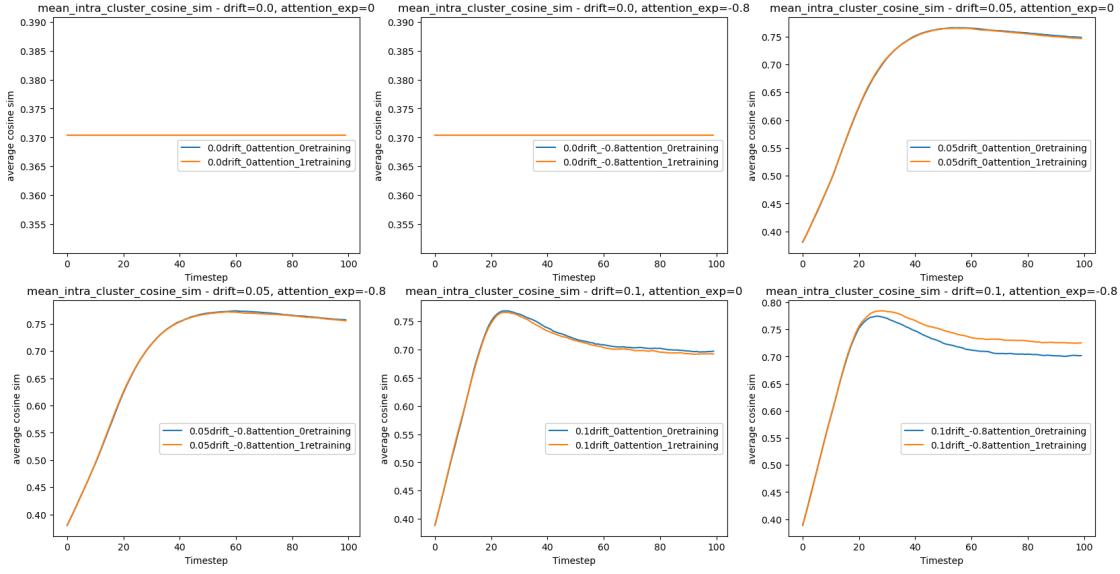
```

[ ]: metric_key = "mean_intra_cluster_cosine_sim"

fig, axes = plt.subplots(2, 3, figsize=(20, 10))

for i in range(len(model_key_pairs)):
    curr_ax = axes[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_key, model_key_pairs[i], ↴
        id_to_readable, mult_sd=0)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"{metric_key} - ↴
        drift={models[model_key_pairs[i][0]][0]}, ↴
        attention_exp={models[model_key_pairs[i][0]][1]}")

```



3.0.2 Graphing Inter -Cluster Cosine Similarity

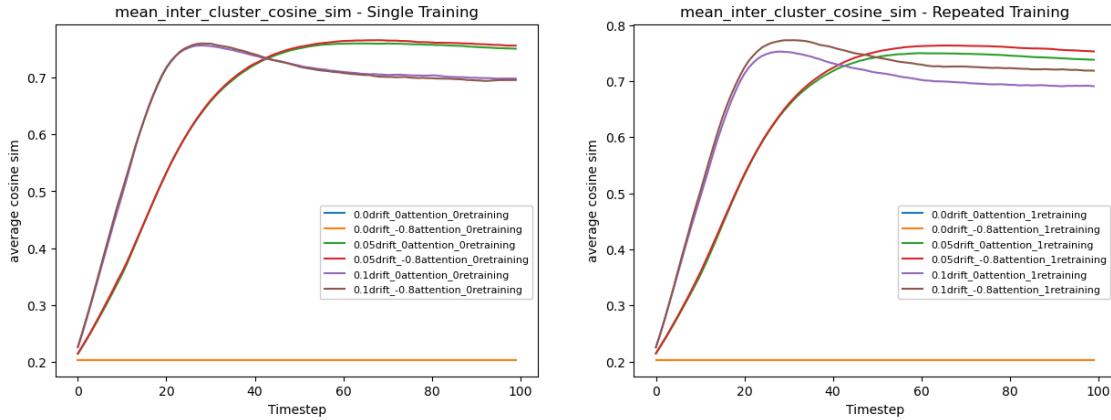
```
[ ]: metric_key = "mean_inter_cluster_cosine_sim"

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

graph_metrics_by_axis(axs[0], results, metric_key, single_training_keys,
    ↪id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right',
    ↪bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axs[1], results, metric_key, repeated_training_keys,
    ↪id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right',
    ↪bbox_to_anchor=(1, 0.5), fontsize="8",)
```

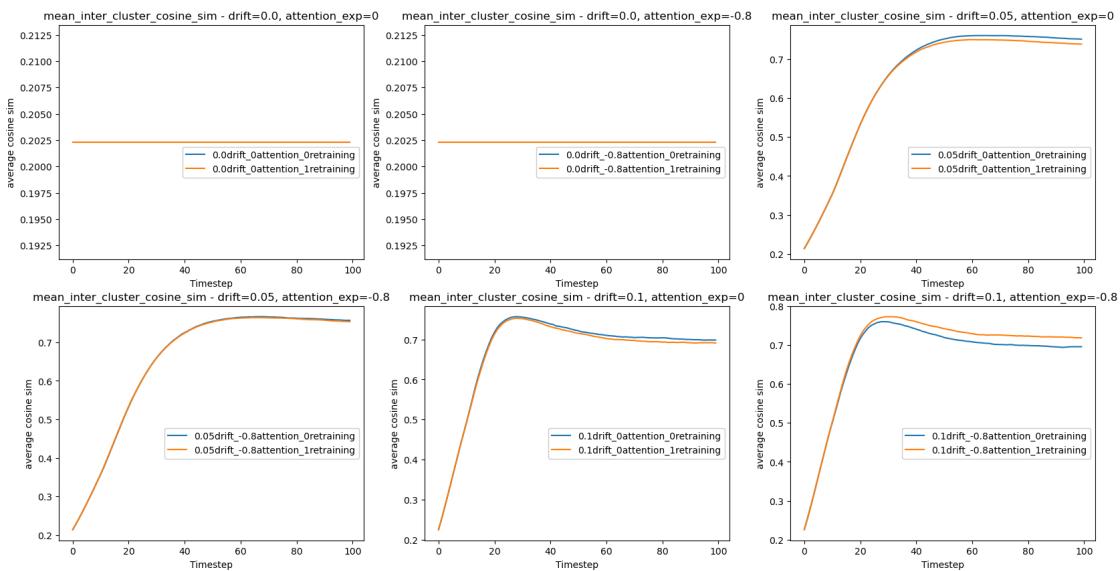
```
[ ]: <matplotlib.legend.Legend at 0x16c36d460>
```



```
[ ]: metric_key = "mean_inter_cluster_cosine_sim"

fig, axs = plt.subplots(2, 3, figsize=(20, 10))

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_key, model_key_pairs[i], ↵
    ↵id_to_readable, mult_sd=0)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"{metric_key} -" ↵
    ↵drift={models[model_key_pairs[i][0]][0]}, ↵
    ↵attention_exp={models[model_key_pairs[i][0]][1]})
```



Graphing Cluster Cosine Similarity relative to Global Cosine Simarilty

```
[ ]: """
Single training
"""

metric_keys = ["mean_intra_cluster_cosine_sim",,
                "mean_inter_cluster_cosine_sim", "mean_global_cosine_sim"]

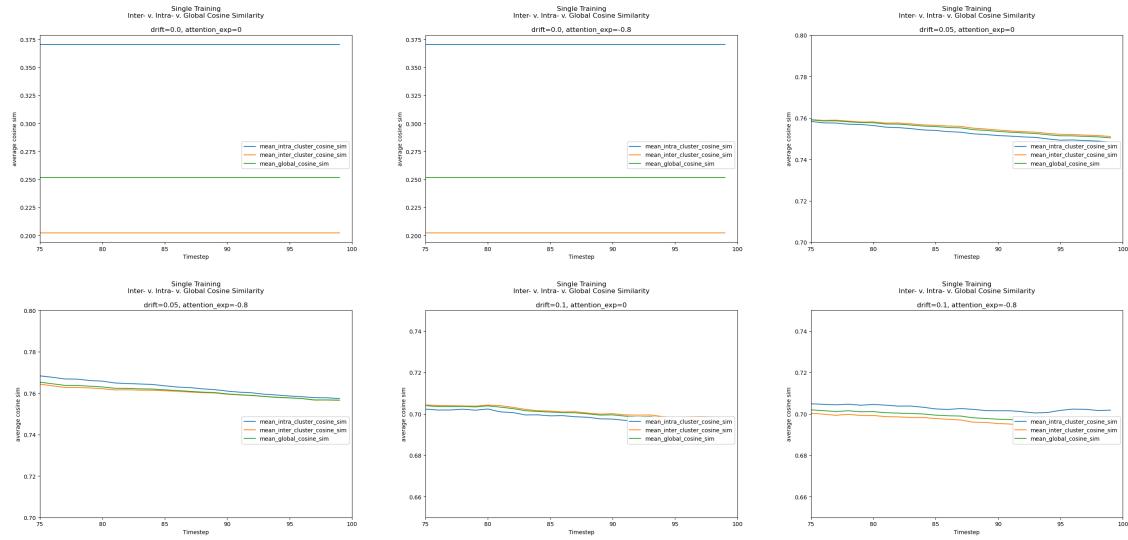
fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    print((model_key_pairs[i][0],))
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_keys[0],,
                           (model_key_pairs[i][0],), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[1],,
                           (model_key_pairs[i][0],), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[2],,
                           (model_key_pairs[i][0],), id_to_readable, mult_sd=0, graph_by="metric")

    curr_ax.set_xlim(75, 100)
    if i > 3:
        curr_ax.set_ylim(0.65, 0.75)
    elif i > 1:
        curr_ax.set_ylim(0.7, 0.8)
    # if i > 1:
    #     curr_ax.set_ylim(0.6, 0.8)
    curr_ax.set_ylabel(y_labels[metric_keys[0]])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Single Training\nInter- v. Intra- v. Global Cosine
Similarity\n\nndrift={models[model_key_pairs[i][0]][0]},",
                      attention_exp={models[model_key_pairs[i][0]][1]})

# fig.tight_layout()

('0.0drift_0attention_0retraining',)
('0.0drift_-0.8attention_0retraining',)
('0.05drift_0attention_0retraining',)
('0.05drift_-0.8attention_0retraining',)
('0.1drift_0attention_0retraining',)
('0.1drift_-0.8attention_0retraining',)
```



```
[ ]: """
Repeated training
"""

metric_keys = ["mean_intra_cluster_cosine_sim", "mean_inter_cluster_cosine_sim", "mean_global_cosine_sim"]

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    print(model_key_pairs[i][1])
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_keys[0], (model_key_pairs[i][1],), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[1], (model_key_pairs[i][1],), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[2], (model_key_pairs[i][1],), id_to_readable, mult_sd=0, graph_by="metric")

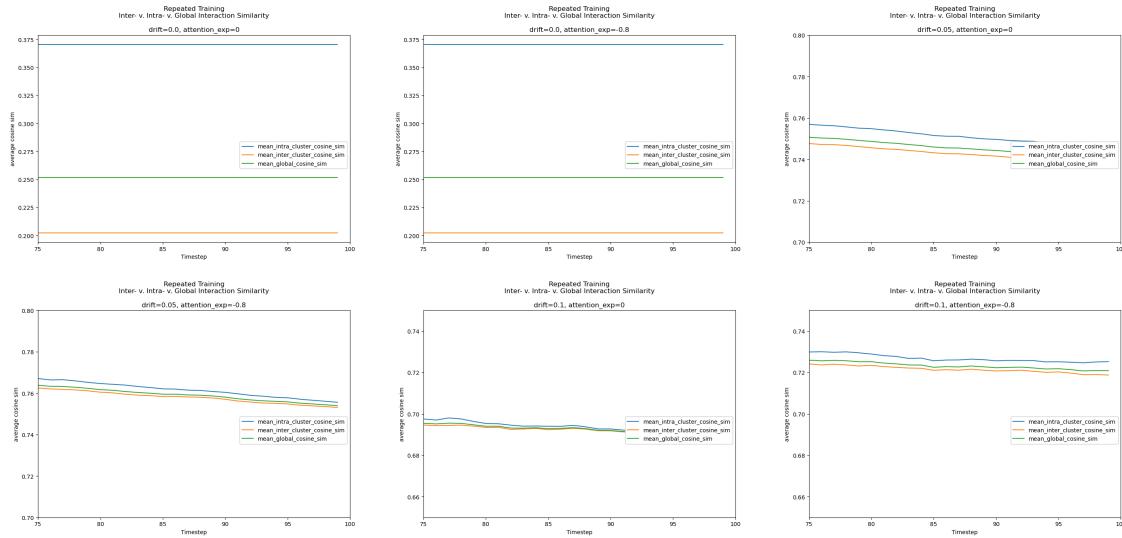
    curr_ax.set_xlim(75, 100)
    if i > 3:
        curr_ax.set_ylim(0.65, 0.75)
    elif i > 1:
        curr_ax.set_ylim(0.7, 0.8)
    # if i > 1:
    #     curr_ax.set_ylim(0.6, 0.8)
    curr_ax.set_ylabel(y_labels[metric_keys[0]])
    curr_ax.set_xlabel("Timestep")
```

```

curr_ax.set_title(f"Repeated Training\nInter- v. Intra- v. Global\nInteraction Similarity\n\ndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")

```

0.0drift_0attention_1retraining
0.0drift_-0.8attention_1retraining
0.05drift_0attention_1retraining
0.05drift_-0.8attention_1retraining
0.1drift_0attention_1retraining
0.1drift_-0.8attention_1retraining



```

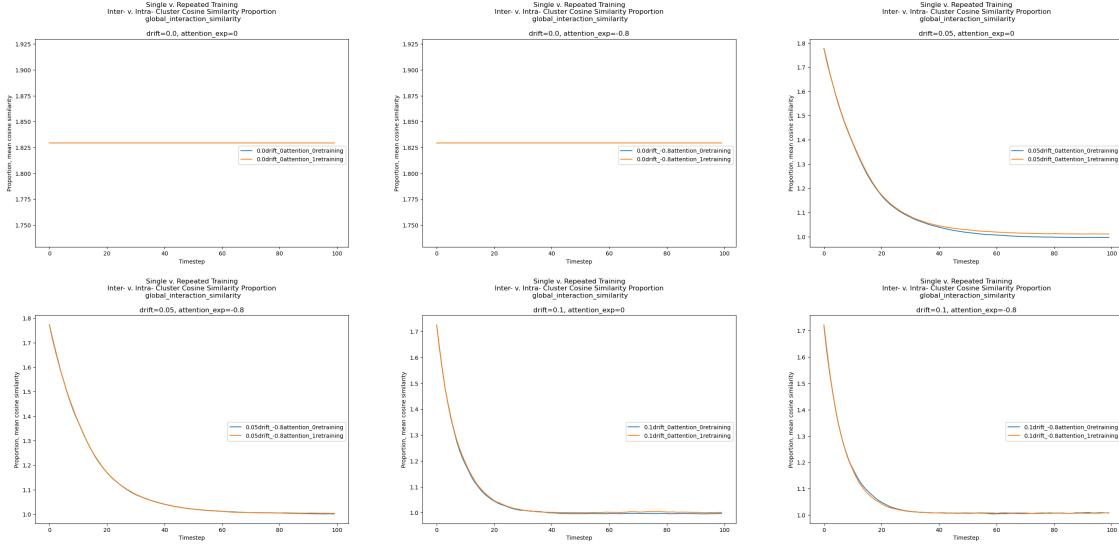
[ ]: numerator = "mean_intra_cluster_cosine_sim"
denominator = "mean_inter_cluster_cosine_sim"

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, denominator, numerator,
                                     model_key_pairs[i], id_to_readable, mult_sd=0)
    curr_ax.set_ylabel("Proportion, mean cosine similarity")
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Single v. Repeated Training\nInter- v. Intra- Cluster\nCosine Similarity Proportion\n\n{global_metric_key}\n\ndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")

    # curr_ax.set_ylim(-0.1, 1.75)
    # curr_ax.set_xlim(40, 100)

```



4 Graphing Distance from centroid

4.0.1 Graphing `mean_cluster_distance_from_centroid`

```
[ ]: metric_key = "mean_cluster_distance_from_centroid"

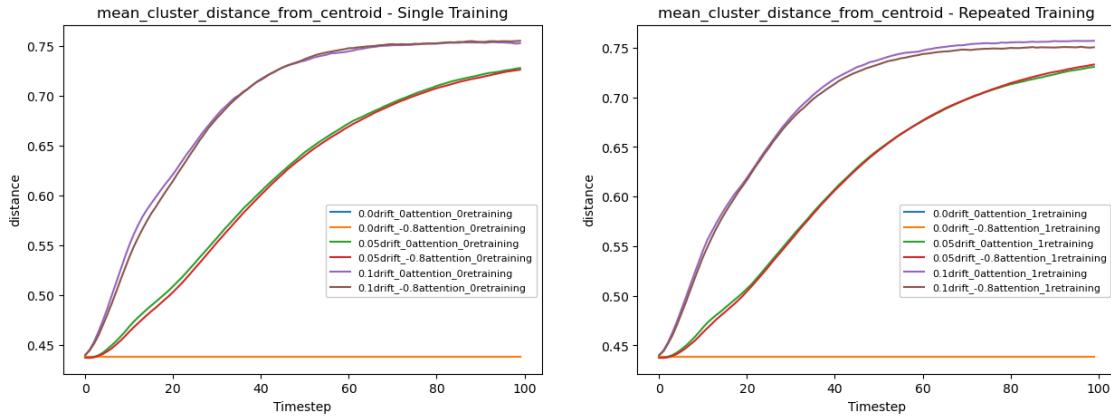
single_training_keys = [key[0] for key in model_key_pairs]
repeated_training_keys = [key[1] for key in model_key_pairs]

fig, axs = plt.subplots(1, 2, figsize=(15, 5))

graph_metrics_by_axis(axs[0], results, metric_key, single_training_keys, ↴
    id_to_readable, mult_sd=0)
axs[0].set_ylabel(y_labels[metric_key])
axs[0].set_xlabel("Timestep")
axs[0].set_title(f"{metric_key} - Single Training")
axs[0].legend(facecolor='white', framealpha=1, loc='upper right', ↴
    bbox_to_anchor=(1, 0.5), fontsize="8",)

graph_metrics_by_axis(axs[1], results, metric_key, repeated_training_keys, ↴
    id_to_readable, mult_sd=0)
axs[1].set_ylabel(y_labels[metric_key])
axs[1].set_xlabel("Timestep")
axs[1].set_title(f"{metric_key} - Repeated Training")
axs[1].legend(facecolor='white', framealpha=1, loc='upper right', ↴
    bbox_to_anchor=(1, 0.5), fontsize="8",)
```

[]: <matplotlib.legend.Legend at 0x177f22e50>



Graphing Intra-Cluster Distance relative to Global Distance

```
[ ]: """
Single training
"""

metric_keys = ['mean_cluster_distance_from_centroid', □
    ↵ 'mean_global_distance_from_centroid']

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    print(model_key_pairs[i][0])
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_keys[0], □
    ↵ (model_key_pairs[i][0],), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[1], □
    ↵ (model_key_pairs[i][0],), id_to_readable, mult_sd=0, graph_by="metric")

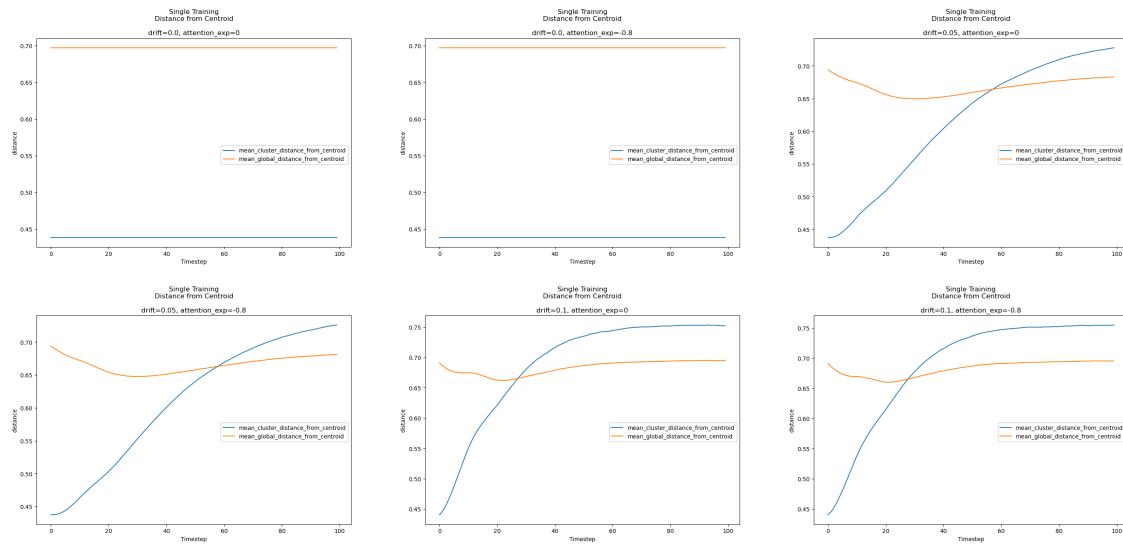
    # curr_ax.set_xlim(75, 100)
    # if i > 3:
    #     curr_ax.set_ylim(0.3, 0.6)
    # elif i > 1:
    #     curr_ax.set_ylim(0.4, 0.8)
    # if i > 1:
    #     curr_ax.set_ylim(0.3, 0.8)
    curr_ax.set_ylabel(y_labels[metric_keys[0]])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Single Training\nDistance from□
    ↵ Centroid\nndrift={models[model_key_pairs[i][0]][0]}, □
    ↵ attention_exp={models[model_key_pairs[i][0]][1]}")
```

0.0drift_0attention_Oretraining

```

0.0drift_-0.8attention_0retraining
0.05drift_0attention_0retraining
0.05drift_-0.8attention_0retraining
0.1drift_0attention_0retraining
0.1drift_-0.8attention_0retraining

```



```

[ ]: """
Repeated training
"""

metric_keys = ['mean_cluster_distance_from_centroid', ↪
    'mean_global_distance_from_centroid']

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    print(model_key_pairs[i][1])
    curr_ax = axs[int(i >= 3), i%3]
    graph_metrics_by_axis(curr_ax, results, metric_keys[0], ↪
        (model_key_pairs[i][1],), id_to_readable, mult_sd=0, graph_by="metric")
    graph_metrics_by_axis(curr_ax, results, metric_keys[1], ↪
        (model_key_pairs[i][1],), id_to_readable, mult_sd=0, graph_by="metric")

    # curr_ax.set_xlim(75, 100)
    # if i > 3:
    #     curr_ax.set_ylim(0.3, 0.6)
    # elif i > 1:
    #     curr_ax.set_ylim(0.4, 0.8)
    # if i > 1:

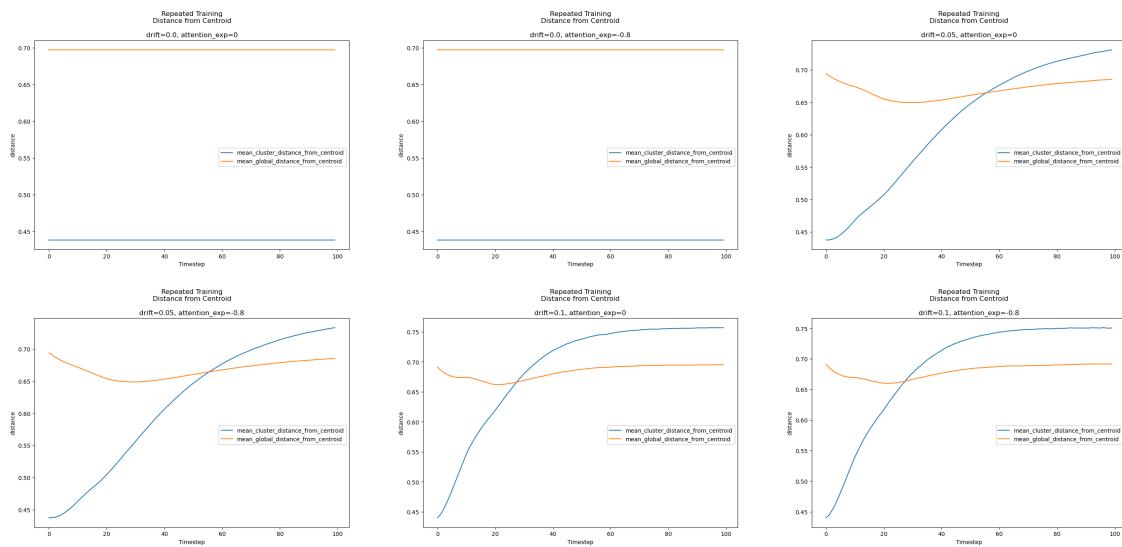
```

```

#     curr_ax.set_ylim(0.3, 0.8)
curr_ax.set_ylabel(y_labels[metric_keys[0]])
curr_ax.set_xlabel("Timestep")
curr_ax.set_title(f"Repeated Training\nDistance from Centroid\nndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")

```

0.0drift_Oattention_1retraining
0.0drift_-0.8attention_1retraining
0.05drift_Oattention_1retraining
0.05drift_-0.8attention_1retraining
0.1drift_Oattention_1retraining
0.1drift_-0.8attention_1retraining



```

[ ]: numerator = "mean_cluster_distance_from_centroid"
denominator = "mean_global_distance_from_centroid"

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

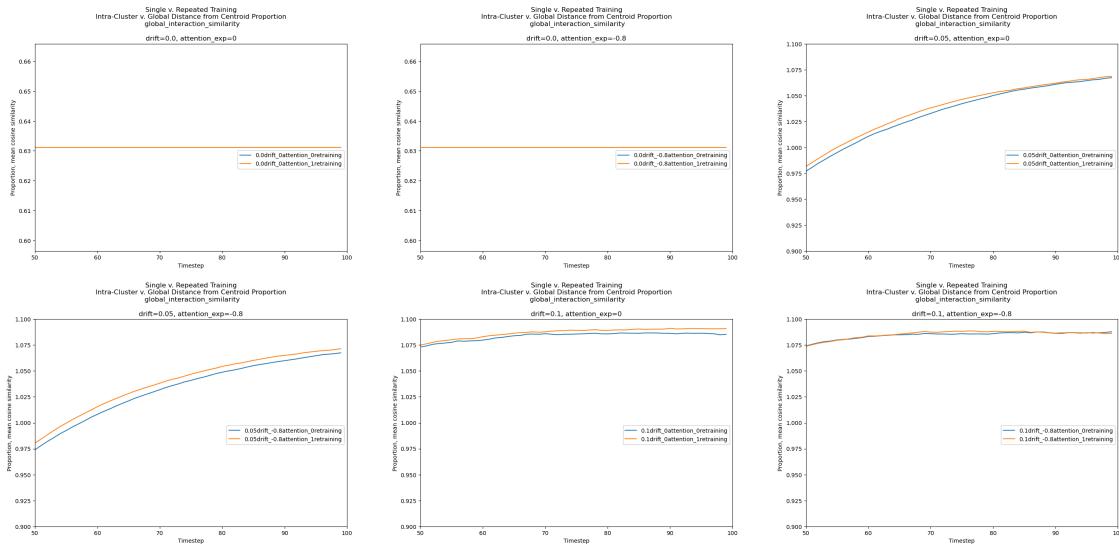
for i in range(len(model_key_pairs)):
    curr_ax = axs[int(i >= 3), i%3]
    graph_relative_to_global_by_axis(curr_ax, results, denominator, numerator,
                                     model_key_pairs[i], id_to_readable, mult_sd=0)
    curr_ax.set_xlim(50, 100)
    if i > 1:
        curr_ax.set_ylim(0.9, 1.1)
    curr_ax.set_ylabel("Proportion, mean cosine similarity")
    curr_ax.set_xlabel("Timestep")

```

```

curr_ax.set_title(f"Single v. Repeated Training\nIntra-Cluster v. Global\nDistance from Centroid Proportion\n{n}\n{global_metric_key}\ndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")
# curr_ax.set_ylim(-0.1, 1.75)
# curr_ax.set_xlim(40, 100)

```



[]:

5 Analyzing histogram data

Graphing `mean_cosine_sim_per_cluster`

```

[ ]: n_clusters = 10
label_map = dict()
for i in range(n_clusters):
    label_map[i] = f"clust_{i}"

```

```

[ ]: """
Single training
"""

metric_key = 'mean_cosine_sim_per_cluster'

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    print(model_key_pairs[i][0])
    curr_ax = axs[int(i >= 3), i%3]

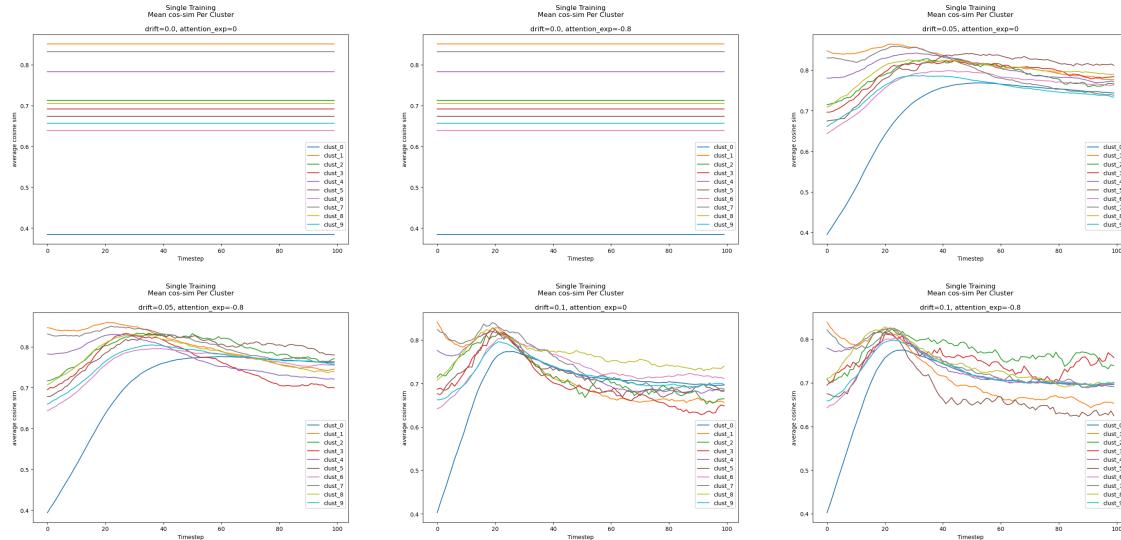
```

```

graph_histogram_metric_by_axis(curr_ax, results, metric_key, u
˓→(model_key_pairs[i][0],), label_map, mean_sigma=0, mult_sd=0, conf_sigma=0)
curr_ax.set_ylabel(y_labels[metric_key])
curr_ax.set_xlabel("Timestep")
curr_ax.set_title(f"Single Training\\nMean cos-sim Per Cluster\\n\\ndrift={models[model_key_pairs[i][0]][0]},\\nattention_exp={models[model_key_pairs[i][0]][1]}")

```

0.0drift_0attention_0retraining
0.0drift_-0.8attention_0retraining
0.05drift_0attention_0retraining
0.05drift_-0.8attention_0retraining
0.1drift_0attention_0retraining
0.1drift_-0.8attention_0retraining



```

[ ]: """
Repeated training
"""

metric_key = 'mean_cosine_sim_per_cluster'

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    print(model_key_pairs[i][1])
    curr_ax = axs[int(i >= 3), i%3]
    graph_histogram_metric_by_axis(curr_ax, results, metric_key, u
˓→(model_key_pairs[i][1],), label_map, mean_sigma=0, mult_sd=0, conf_sigma=0)
    curr_ax.set_ylabel(y_labels[metric_key])

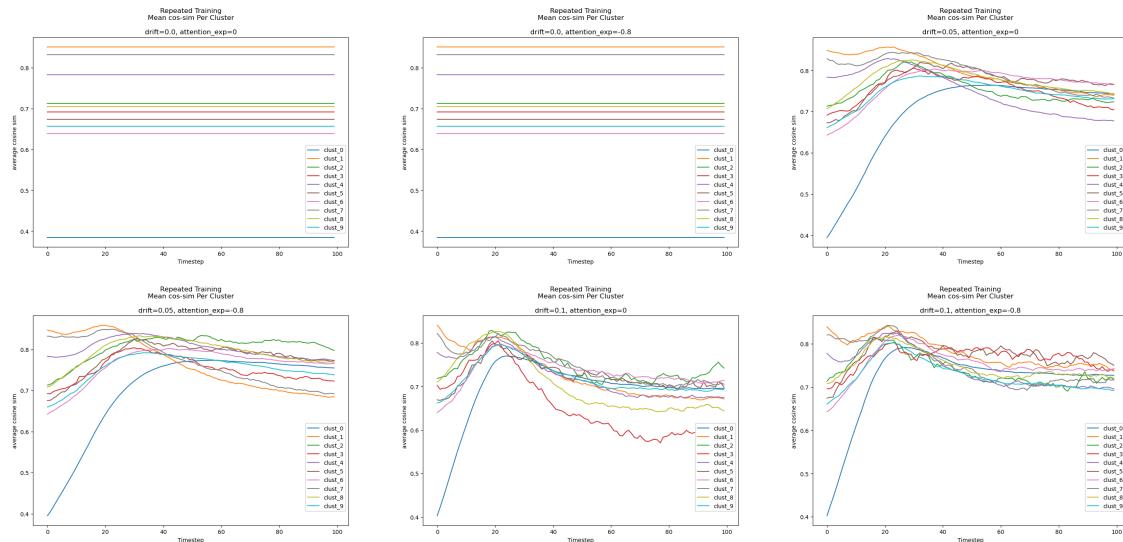
```

```

curr_ax.set_xlabel("Timestep")
curr_ax.set_title(f"Repeated Training\nMean cos-sim Per Cluster\nndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")

```

0.0drift_0attention_1retraining
 0.0drift_-0.8attention_1retraining
 0.05drift_0attention_1retraining
 0.05drift_-0.8attention_1retraining
 0.1drift_0attention_1retraining
 0.1drift_-0.8attention_1retraining



Graphing `mean_distance_from_centroid_per_cluster`

```

[ ]:
"""
Single training
"""

metric_key = 'mean_distance_from_centroid_per_cluster'

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):
    #
    #print(results['mean_global_distance_from_centroid'][model_key_pairs[i][0]].mean(axis=0).shape)
    # break
    print(model_key_pairs[i][0])
    curr_ax = axs[int(i >= 3), i%3]

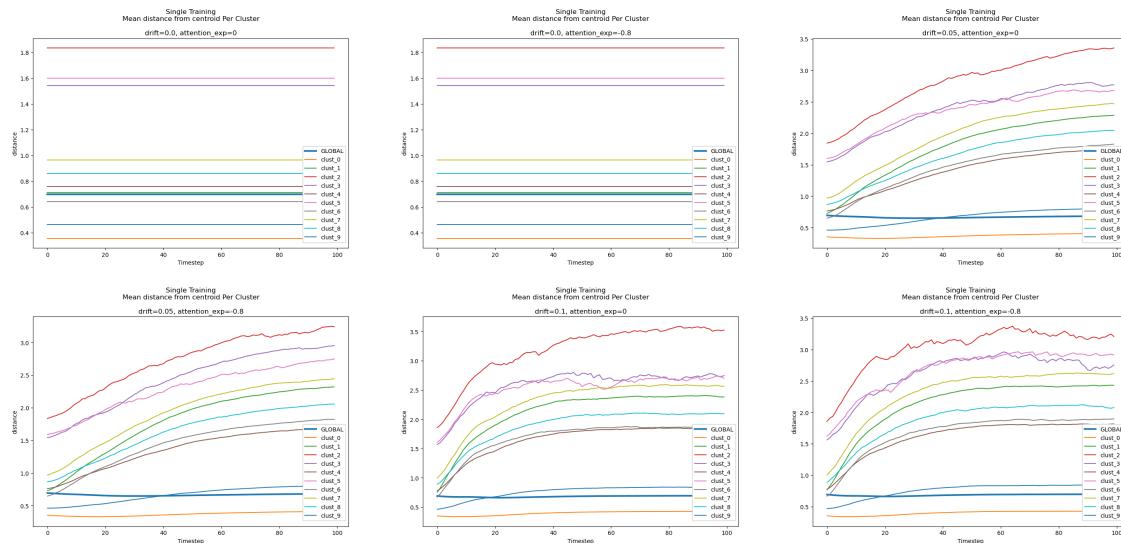
```

```

    curr_ax.
    ↪plot(results['mean_global_distance_from_centroid'][model_key_pairs[i][0]].
    ↪mean(axis=0), label="GLOBAL", linewidth=3)
    graph_histogram_metric_by_axis(curr_ax, results, metric_key,
    ↪(model_key_pairs[i][0],), label_map, mean_sigma=0, mult_sd=0, conf_sigma=0)
    curr_ax.set_ylabel(y_labels[metric_key])
    curr_ax.set_xlabel("Timestep")
    curr_ax.set_title(f"Single Training\nMean distance from centroid Per Cluster\nndrift={models[model_key_pairs[i][0]][0]},\nattention_exp={models[model_key_pairs[i][0]][1]}")

```

0.0drift_0attention_0retraining
0.0drift_-0.8attention_0retraining
0.05drift_0attention_0retraining
0.05drift_-0.8attention_0retraining
0.1drift_0attention_0retraining
0.1drift_-0.8attention_0retraining



[]:

```

"""
Repeated training
"""

metric_key = 'mean_distance_from_centroid_per_cluster'

fig, axs = plt.subplots(2, 3, figsize=(30, 15))
fig.tight_layout(pad=10.0)

for i in range(len(model_key_pairs)):

```

```

#_
print(results['mean_global_distance_from_centroid'][model_key_pairs[i][0]].
      mean(axis=0).shape)

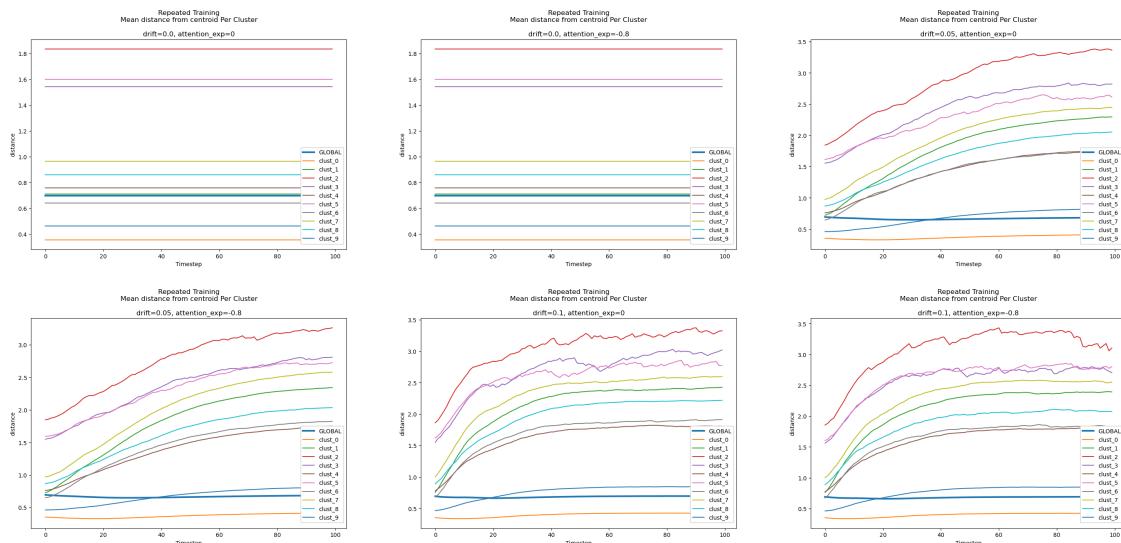
# break
print(model_key_pairs[i][1])
curr_ax = axs[int(i >= 3), i%3]
curr_ax.

plot(results['mean_global_distance_from_centroid'][model_key_pairs[i][1]].
      mean(axis=0), label="GLOBAL", linewidth=3)

graph_histogram_metric_by_axis(curr_ax, results, metric_key,_
      (model_key_pairs[i][1],), label_map, mean_sigma=0, mult_sd=0, conf_sigma=0)
curr_ax.set_ylabel(y_labels[metric_key])
curr_ax.set_xlabel("Timestep")
curr_ax.set_title(f"Repeated Training\nMean distance from centroid Per Cluster\nndrift={models[model_key_pairs[i][0]][0]},"
      attention_exp={models[model_key_pairs[i][0]][1]})


```

0.0drift_0attention_1retraining
 0.0drift_-0.8attention_1retraining
 0.05drift_0attention_1retraining
 0.05drift_-0.8attention_1retraining
 0.1drift_0attention_1retraining
 0.1drift_-0.8attention_1retraining



[]: results.keys()

[]: dict_keys(['mse', 'interaction_spread', 'global_interaction_similarity',
 'inter_cluster_interaction_similarity', 'intra_cluster_interaction_similarity',
 'mean_global_cosine_sim', 'mean_intra_cluster_cosine_sim',

```
'mean_inter_cluster_cosine_sim', 'mean_cosine_sim_per_cluster',
'mean_cluster_distance_from_centroid', 'mean_global_distance_from_centroid',
'mean_distance_from_centroid_per_cluster'])
```

```
[ ]:
```